

Linear Supervised Learning

M. Ndaoud



Machine Learning - Recap

Supervised Learning

Definition: Supervised learning uses labeled data to train models. The model learns from input-output pairs and can predict outcomes for new, unseen data.

Common Use Cases: Classification (spam detection, disease diagnosis) and Regression (predicting housing prices, stock trends).

Key Methods: Linear Regression, Logistic Regression, Decision Trees, Random Forests, Support Vector Machines (SVM), K-Nearest Neighbors (KNN), Neural Networks

Unsupervised Learning

Definition: Unsupervised learning works with unlabeled data. The model finds patterns, structures, or groupings without pre-existing labels or categories.

Common Use Cases: Clustering (customer segmentation, image grouping) and Dimensionality Reduction (reducing features while preserving information).

Key Methods: K-Means, Clustering, Hierarchical Clustering, Principal Component Analysis (PCA), DBSCAN (Density-Based Spatial Clustering), Association Rule Learning (e.g., Apriori, used for market basket analysis)

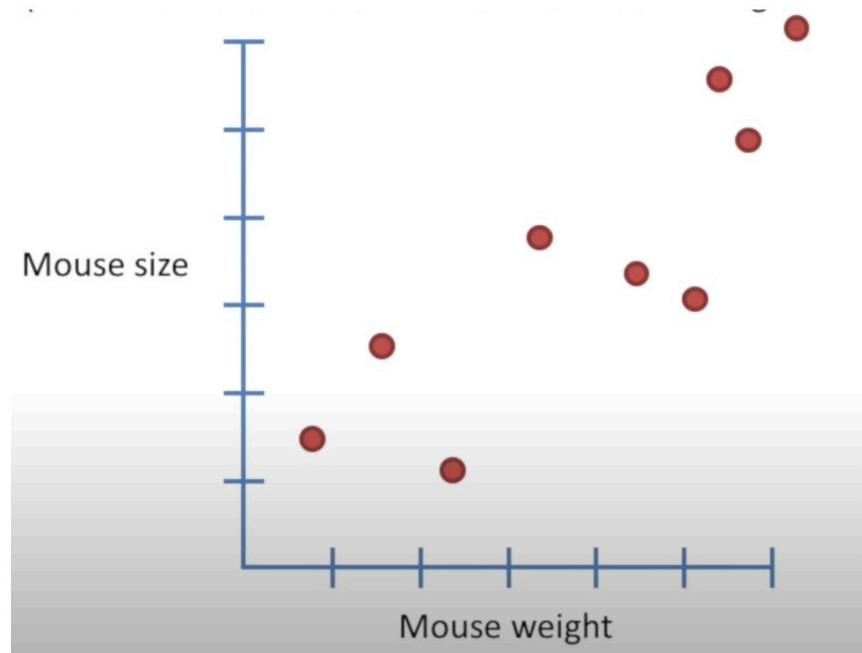
Linear Regression - Recap

Predicting continuous outcomes

Can we predict the size of a mouse based on its weight?

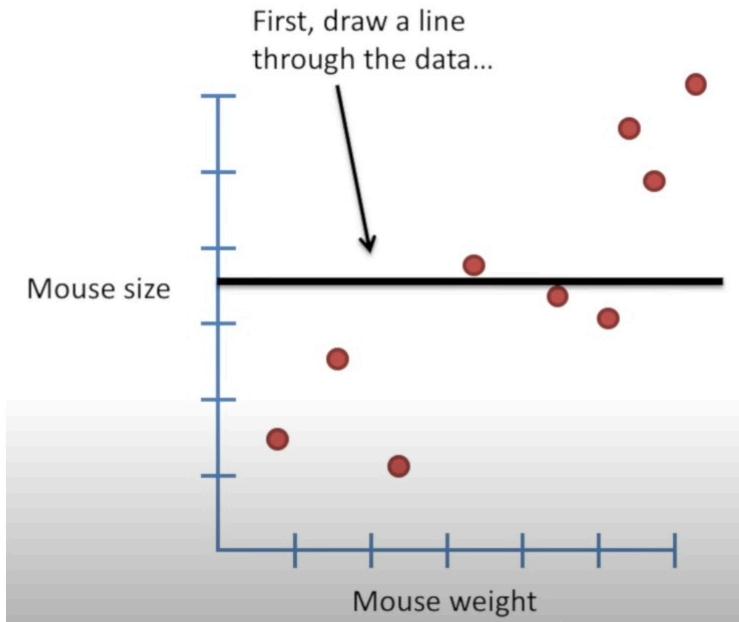
Here are a few data points we have collected on mouses.

Linear regression is about fitting a line to these data.



Linear Regression - Recap

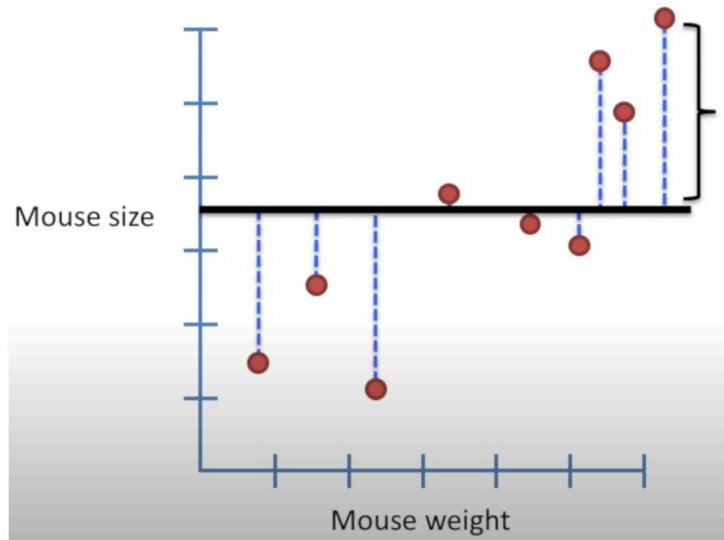
Predicting continuous outcomes



First, draw a line through the data

Linear Regression - Recap

Predicting continuous outcomes



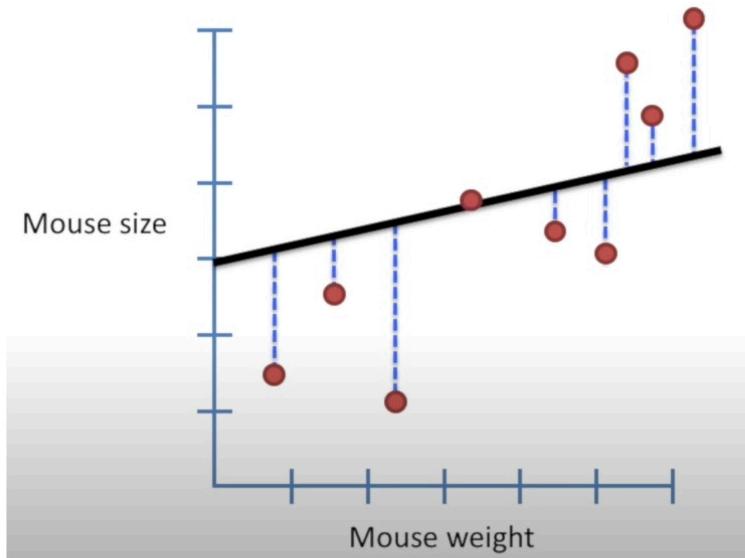
Second, calculate sum of squared residuals

Measure the distance from the line to each data point, square these distances, and add them up

Note: We square them so that negative distances do not cancel out positive distances

Linear Regression - Recap

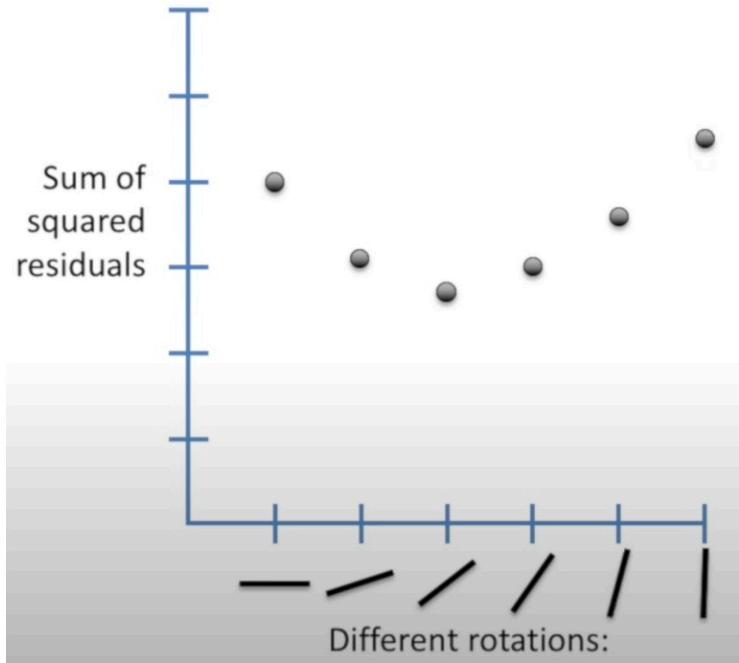
Predicting continuous outcomes



Third, rotate the line and
reproduce the process

Linear Regression - Recap

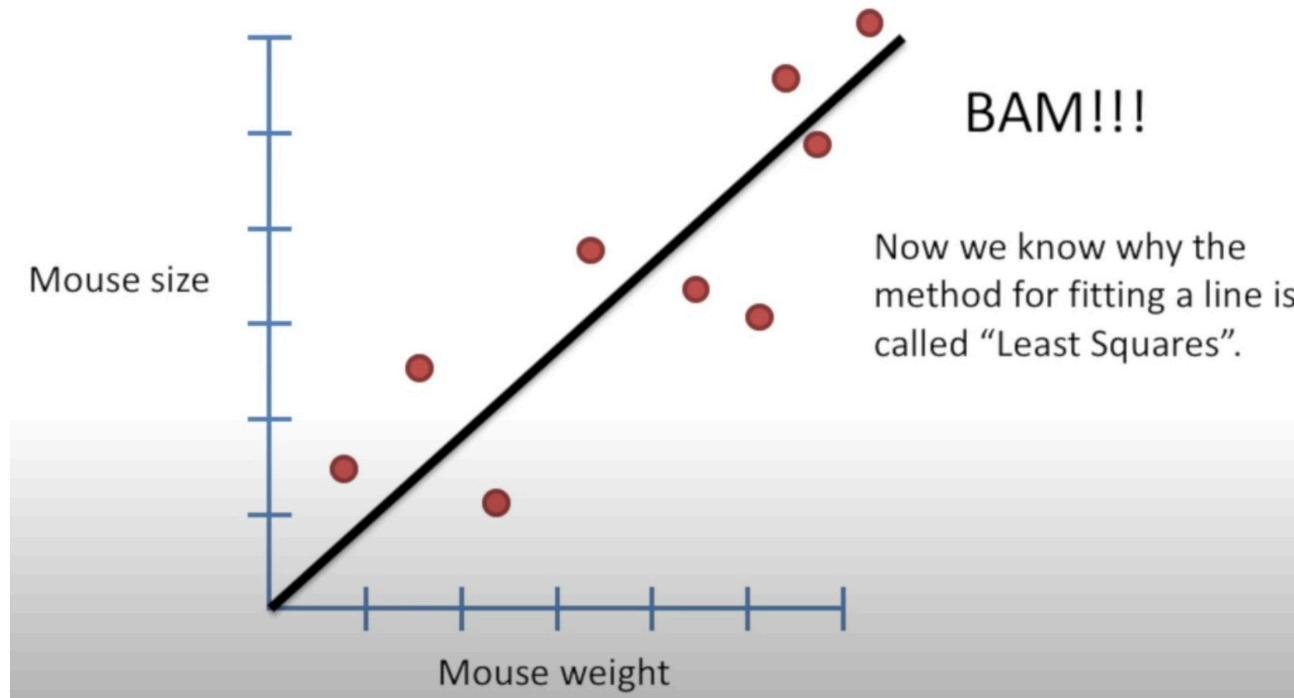
Predicting continuous outcomes



Finally, plot the results and find the rotation that has the least sum of squared residuals

Linear Regression - Recap

Predicting continuous outcomes



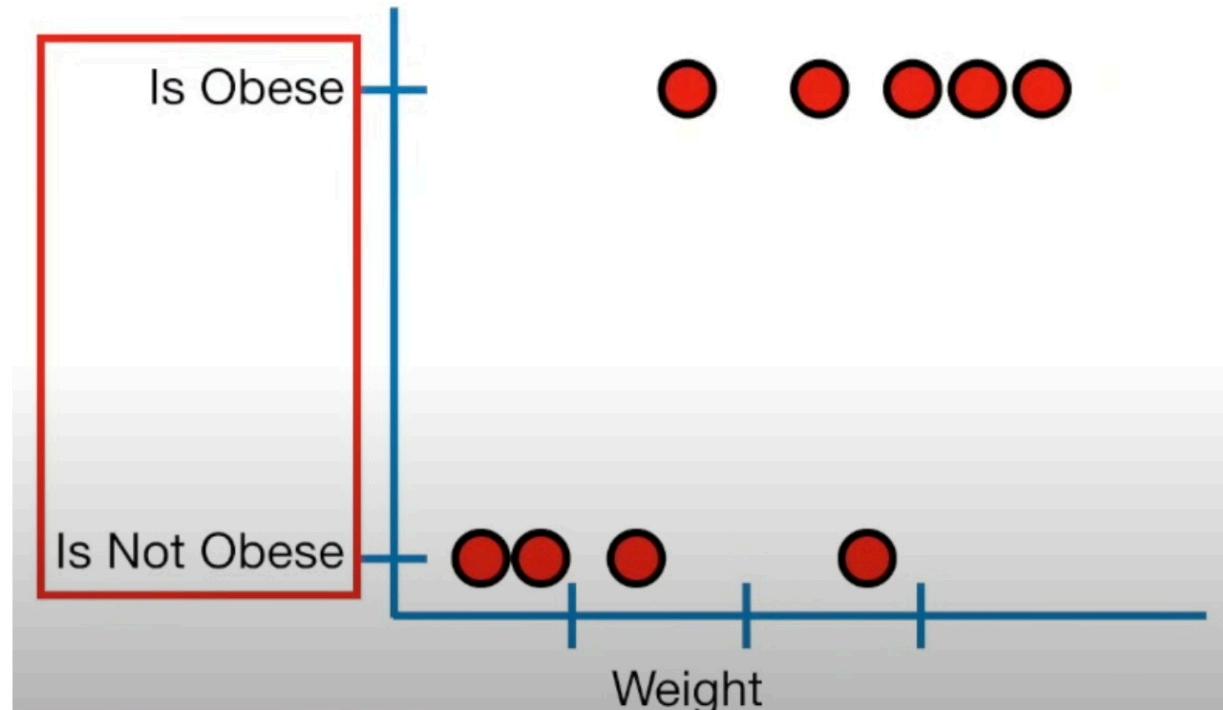
And we can predict the size of any new mouse, based on its weight !

Credits to StatQuest for these slides

This is great, but what if instead of the size, we want to predict if a mouse is obese (Y/N) based on its weight?

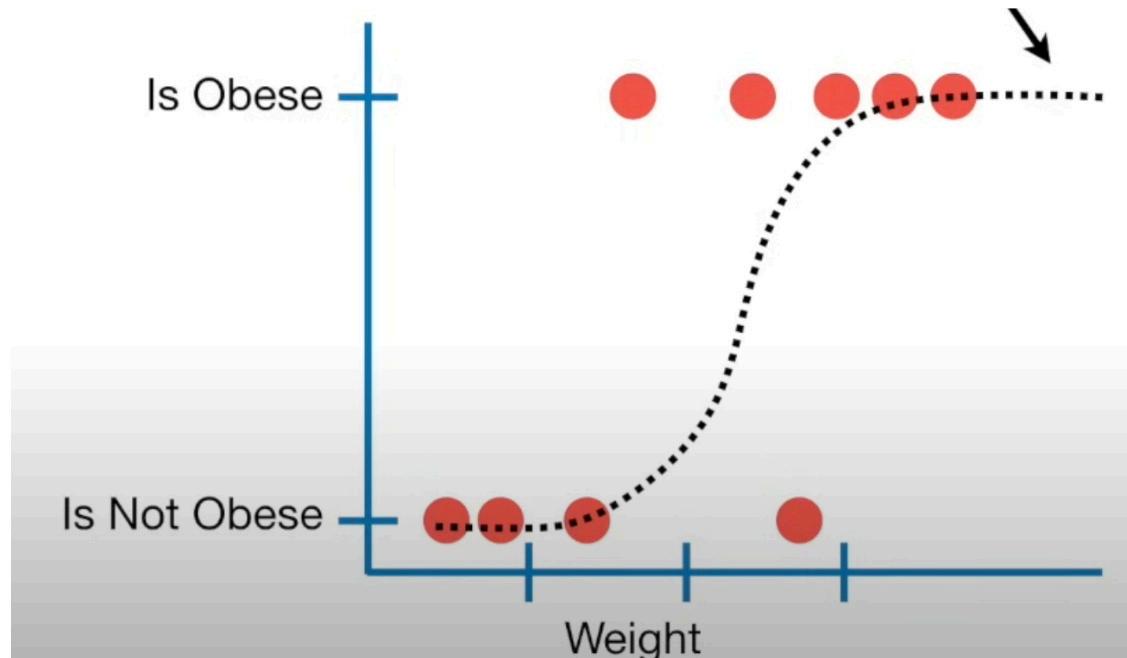
Logistic Regression - Intuition

If you are trying to fit a straight line here, and predict values based on that, you might end up with unrealistic values...

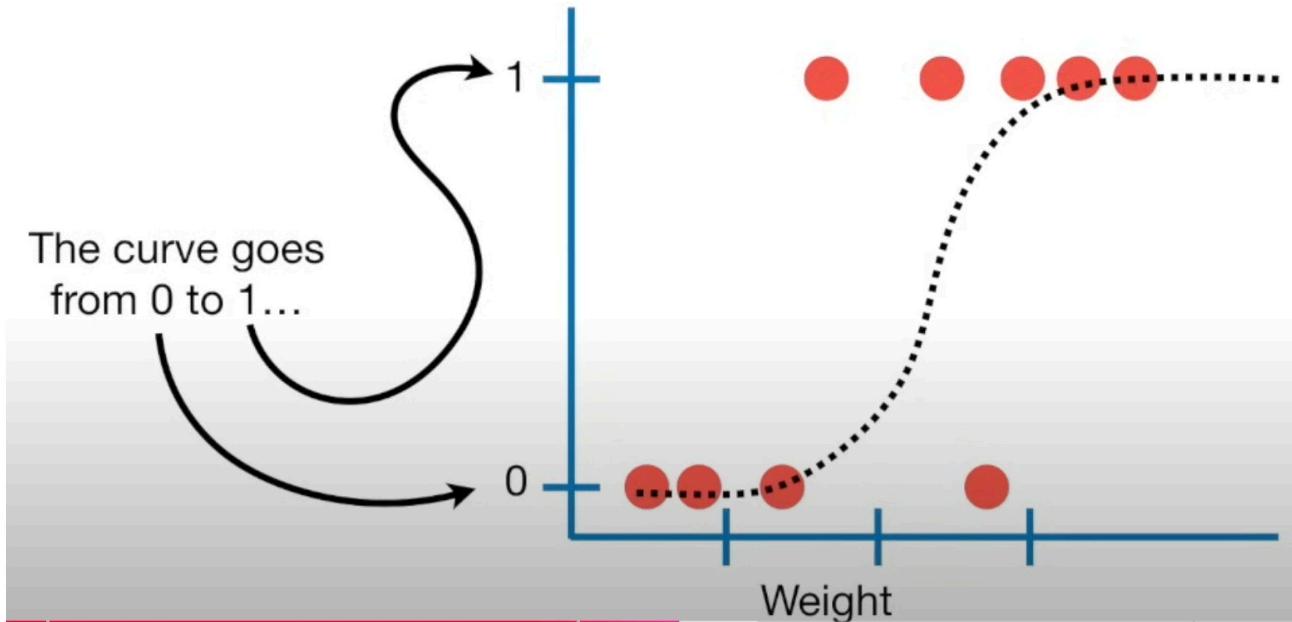


Logistic Regression - Intuition

Instead, we fit an S-curve called "Logistic Function"

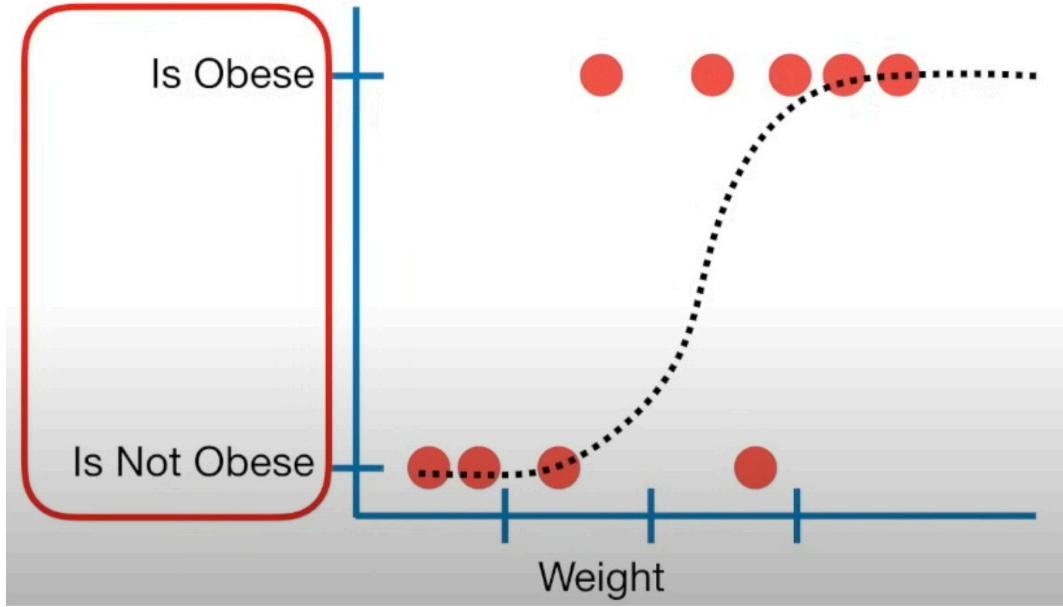


Logistic Regression - Intuition



...and represents the probability that a mouse is obese given its weight

Logistic Regression - Intuition



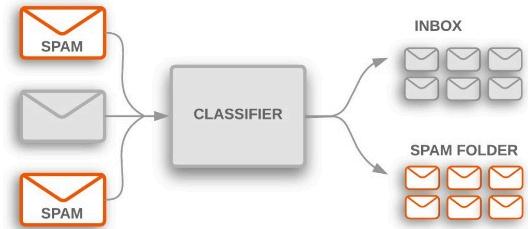
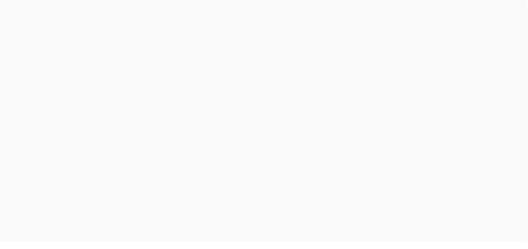
Logistic regression can be used to classify samples...

.. but also to assess what variables are useful for classifying samples

What is Classification?

Predicting the category or class to which a new observation belongs.

Classification models can be divided into two main types: binary classification, where there are only two possible classes, and multi-class classification, where there are multiple classes to choose from.



Classification Tasks

- **Spam detection**

Classifying emails as spam or non-spam to protect users from unwanted messages.

- **Disease prediction**

Identifying the likelihood of a person developing a specific disease based on their medical history and symptoms.

- **Sentiment analysis**

Determining the emotional tone of text, such as reviews or social media posts, to gauge public opinion.

- **Image recognition**

Identifying and classifying the contents of images, such as objects, people, or scenes.

- **Fraud detection**

Identifying suspicious financial transactions or activities to prevent and detect fraud.

Classification Process

- **Data collection and preprocessing**

Gather relevant data, clean and preprocess it to prepare it for feature engineering and model training. This may involve handling missing values, outliers, and encoding categorical variables.

- **Feature engineering**

Identify and create relevant features from the preprocessed data that can help the model learn patterns and make accurate predictions. This may include feature selection, transformation, and dimensionality reduction techniques.

- **Model selection and training**

Select an appropriate machine learning algorithm for the classification task and train the model using the engineered features. This involves hyperparameter tuning and optimization to achieve the best model performance.

- **Model evaluation**

Evaluate the trained model's performance using appropriate metrics, such as accuracy, precision, recall, and F1-score. This step helps assess the model's effectiveness and identify areas for further improvement.

- **Model deployment**

Deploy the trained and evaluated model into a production environment, where it can be used to make real-time predictions on new, unseen data. This may involve integrating the model with a web application or an API.

Popular Classification Algorithms

- **Logistic Regression**

A statistical model used for binary classification tasks, where the goal is to predict whether an input belongs to one of two classes.

- **Decision Trees**

A tree-based algorithm that makes predictions by recursively partitioning the input space based on feature values, creating a set of rules to make decisions.

- **Random Forests**

An ensemble learning method that combines multiple decision trees to improve the accuracy and stability of the predictions, reducing the risk of overfitting.

- **Support Vector Machines (SVMs)**

A supervised learning algorithm that finds the optimal hyperplane to separate different classes, maximizing the margin between the closest data points of each class.

- **Neural Networks**

A machine learning model inspired by the human brain, consisting of interconnected nodes (neurons) that learn to map input data to output predictions through a process of optimization.

Coefficients:

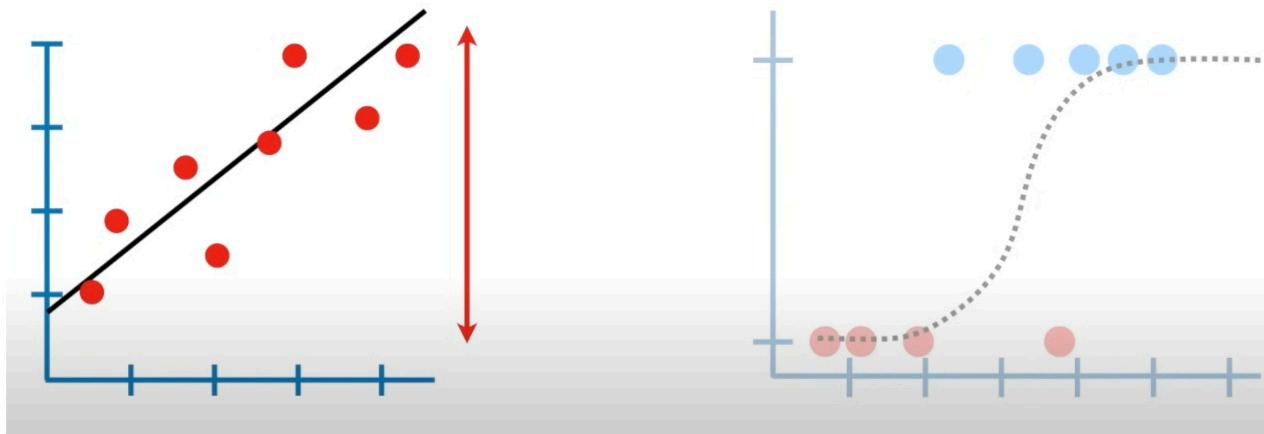
	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.476	2.364	-1.471	0.1414
weight	1.825	1.088	1.678	0.0934

This is what you will get from your logistic regression model when you run it in your computer

I would like you to be able to understand and interpret those coefficients.

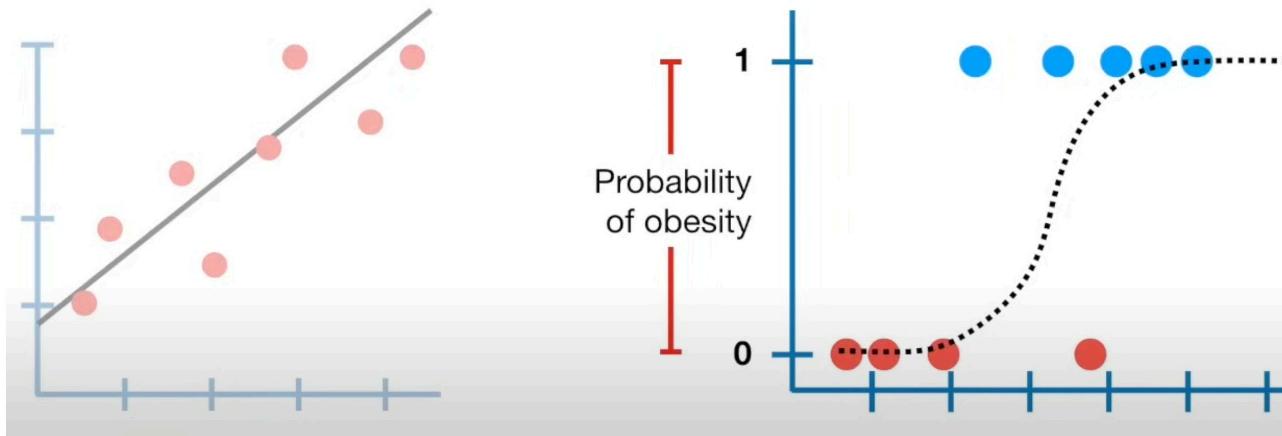
Logistic Regression - Coefficients

With linear regression, the values on the y-axis can, in theory, be any number...

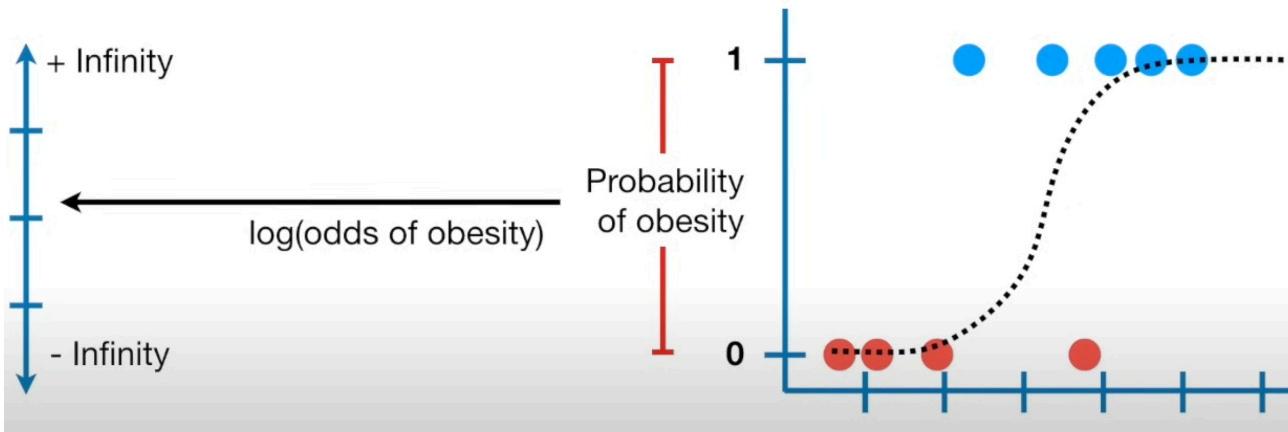


Logistic Regression - Coefficients

...unfortunately, with logistic regression, the y-axis is confined to probability values between 0 and 1.



Logistic Regression - Coefficients

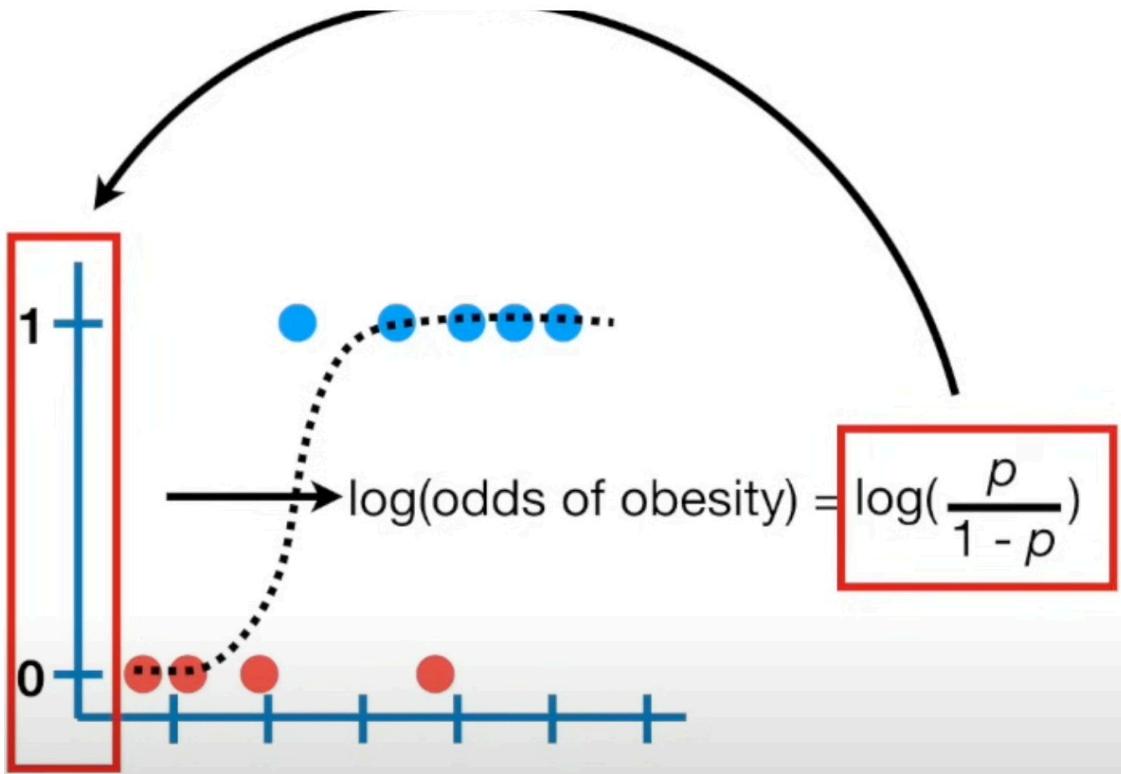


To solve the problem, we transform the probability into the $\log(\text{odds of obesity})$

Odds = probability of an event
happening/probability of an event not happening

Taking the $\log(\text{odd})$ transforms the problem into one where we can use a straight line

Logistic Regression - Coefficients

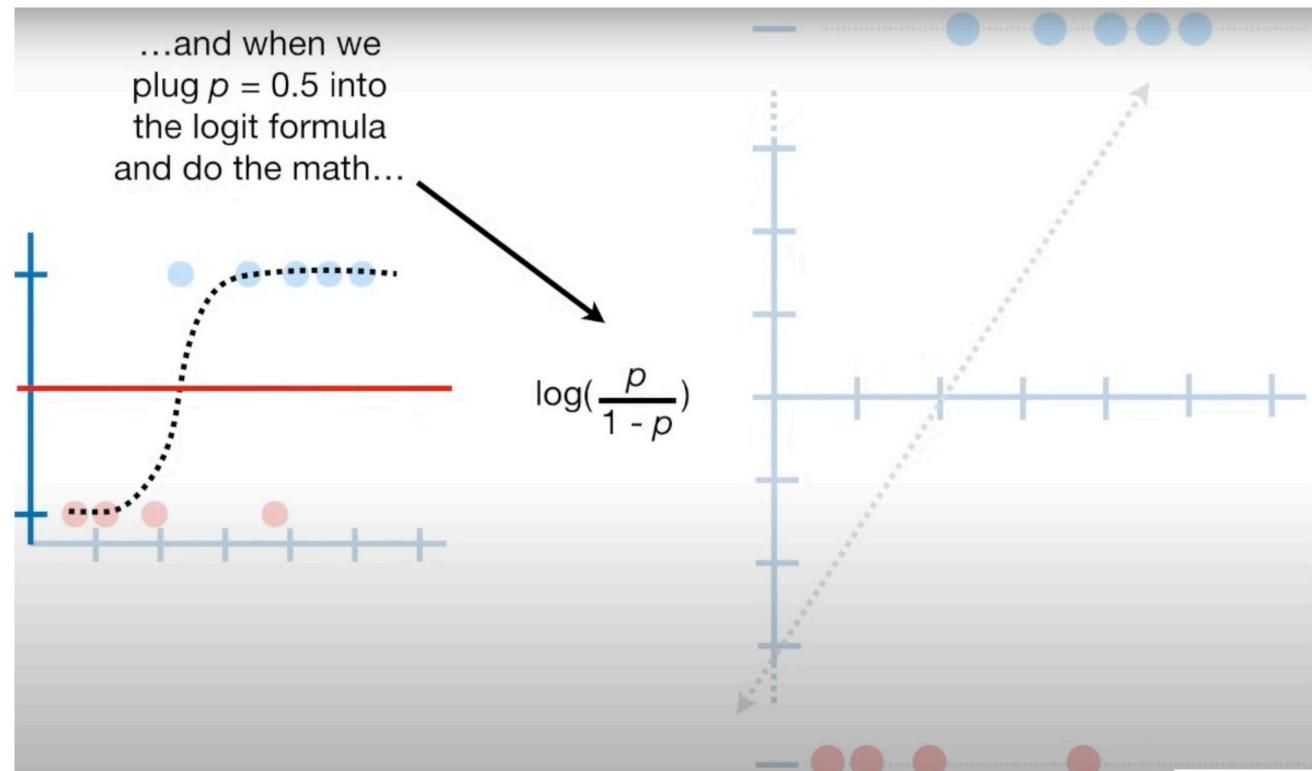


This is the pretty formula...

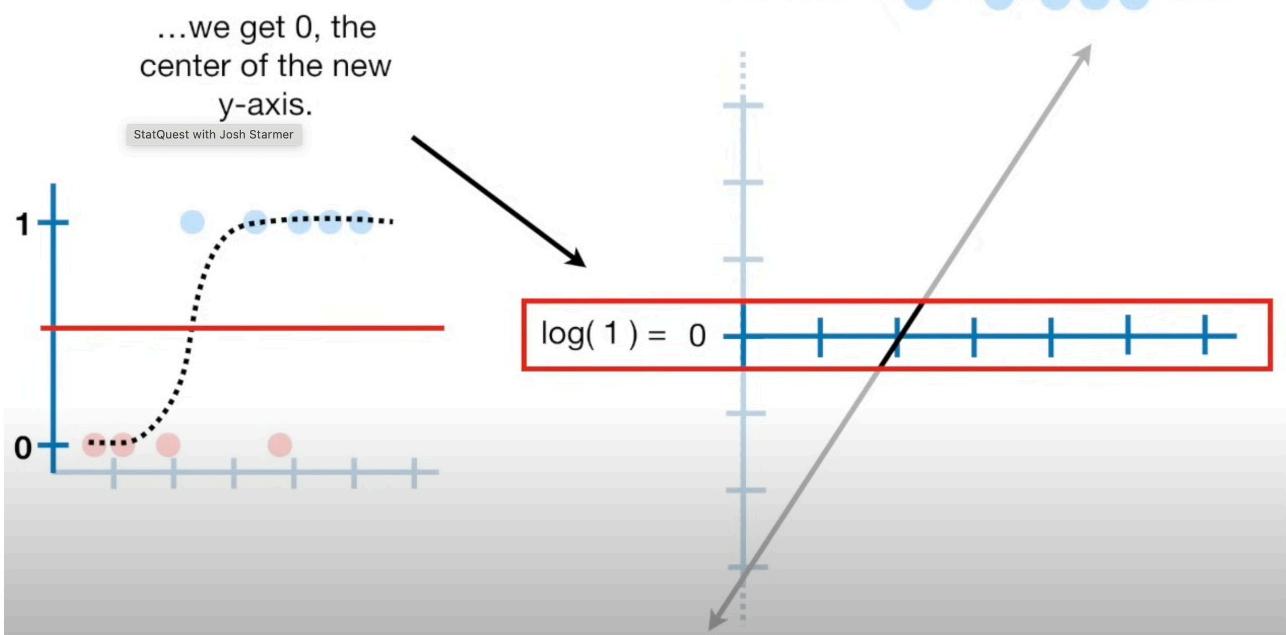
... and let's see how it works !

Logistic Regression - Coefficients

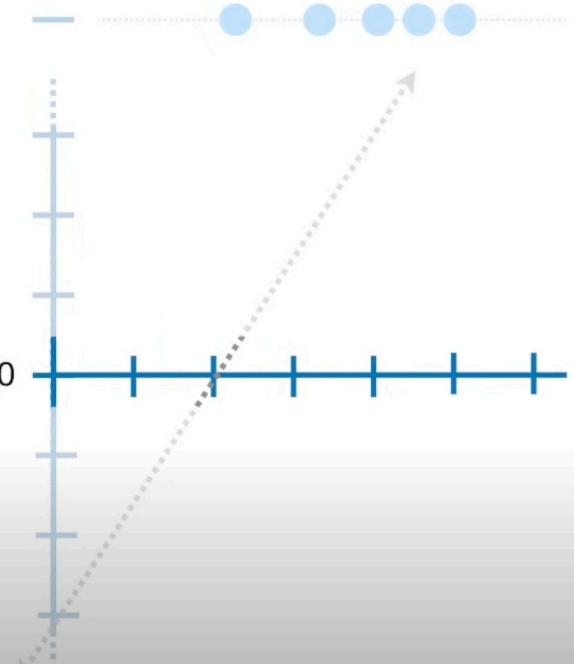
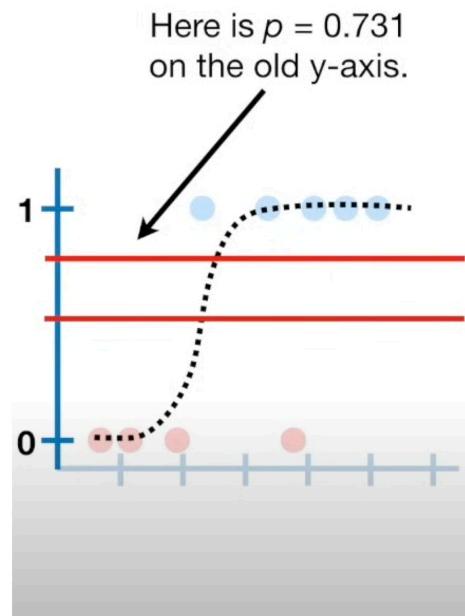
The mid point on the left corresponds to a probability of 0.5...



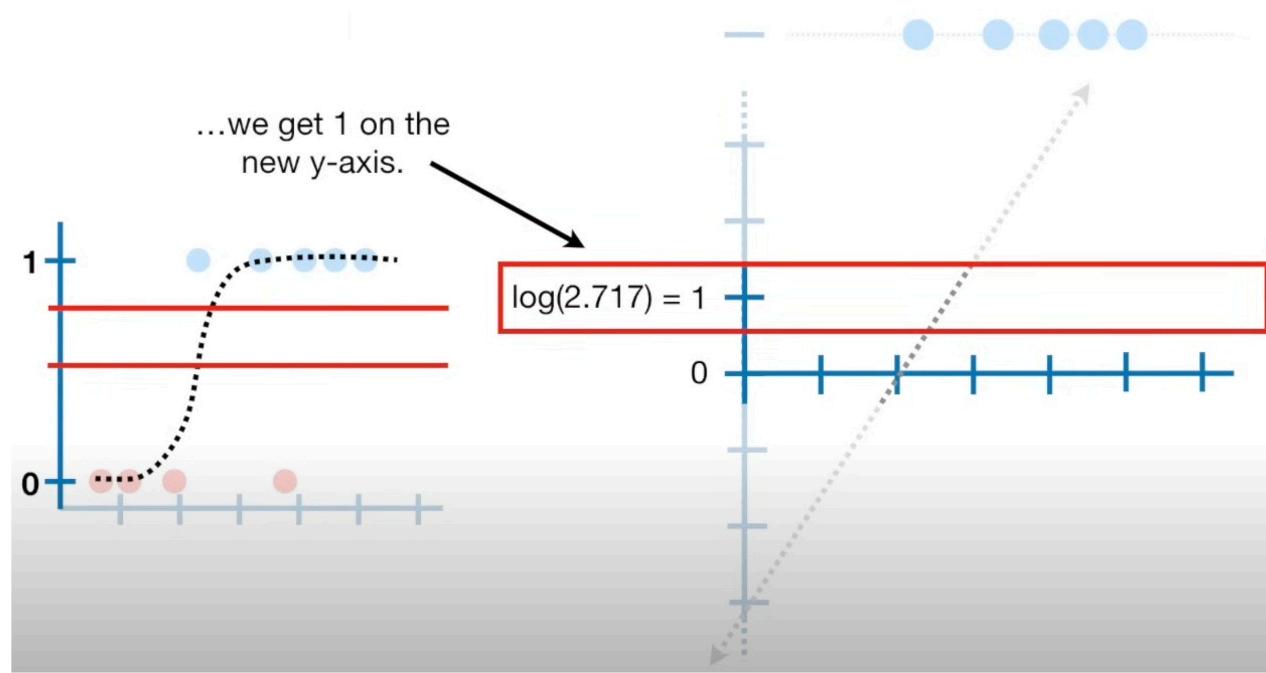
Logistic Regression - Coefficients



Logistic Regression - Coefficients

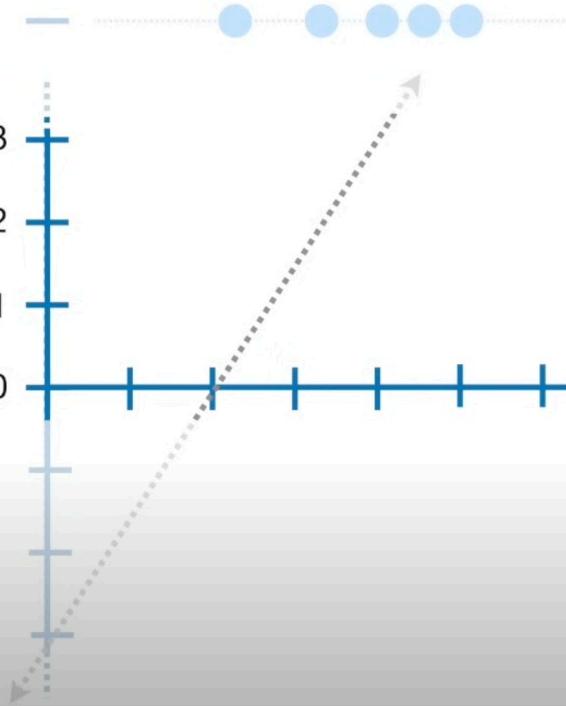
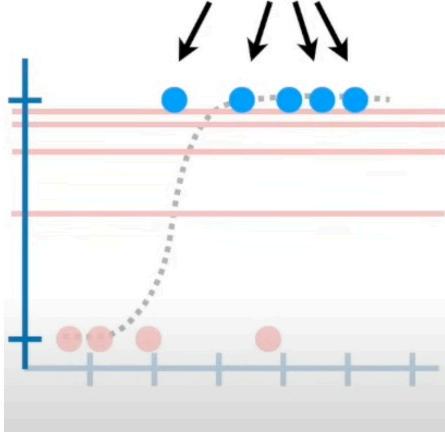


Logistic Regression - Coefficients



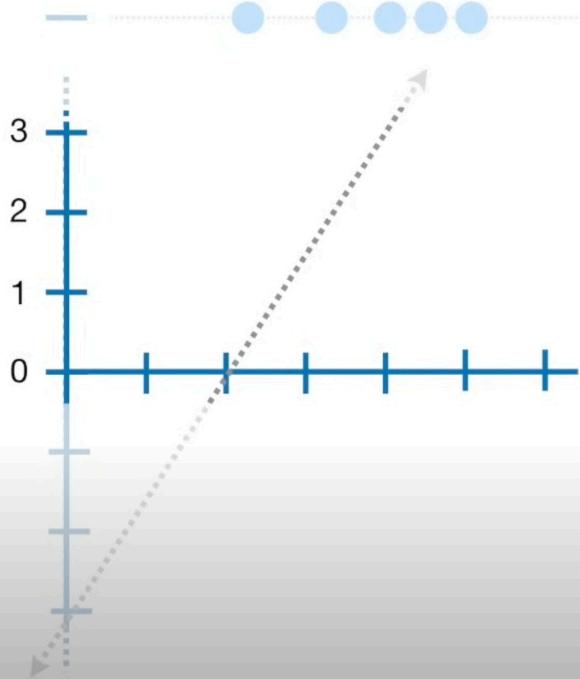
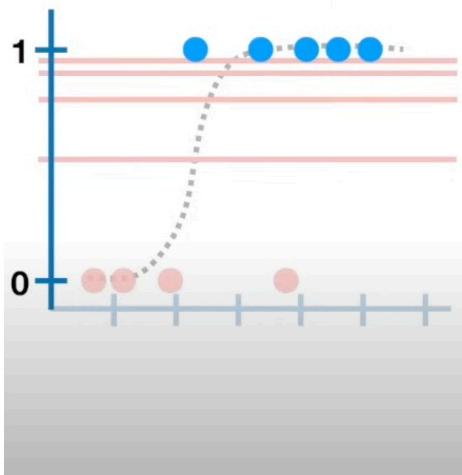
Logistic Regression - Coefficients

Lastly, these blue points from the original data are at $p = 1...$



Logistic Regression - Coefficients

If we plug $p = 1$ into
the logit function and → do the math...



We get $\log(1) - \log(0)$

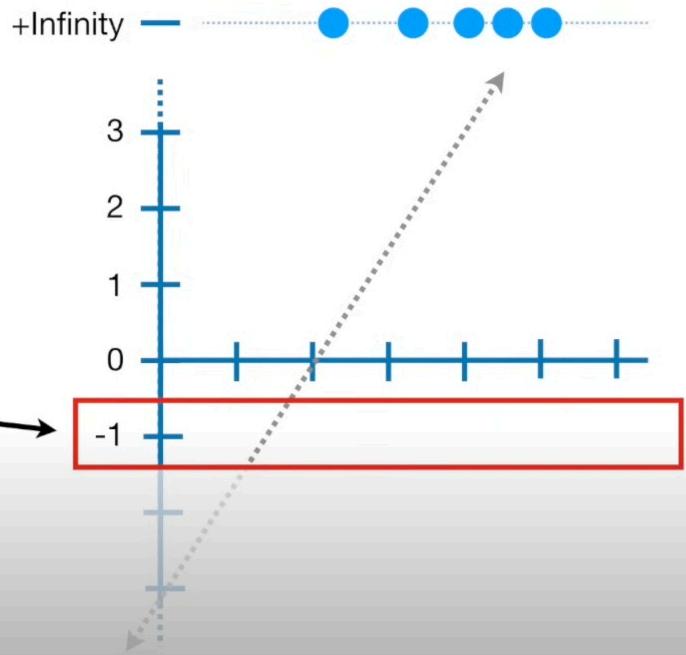
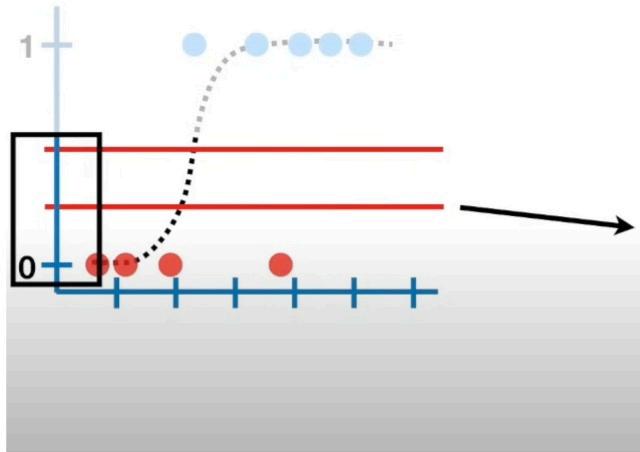
As $\log(0)$ is defined
as $-\infty$

And something $- (-\infty) = +\infty$

This means the
original data points
labeled "obese" are
now at $+\infty$ on the
new Y axis

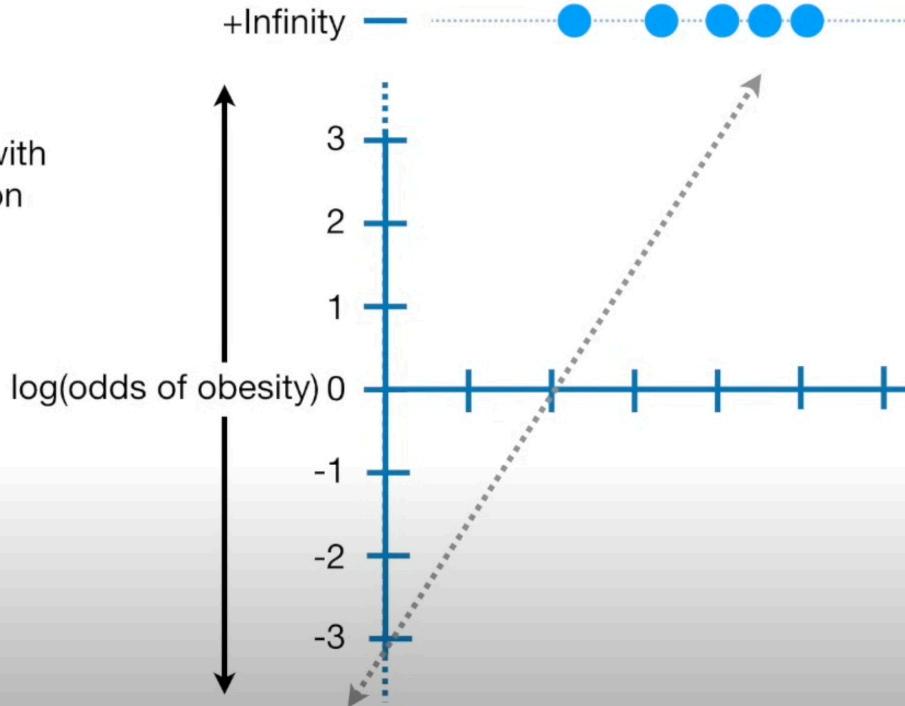
Logistic Regression - Coefficients

Similarly, 0.5 to 0 on the old y-axis is stretched out from 0 to $-\infty$ on the new y-axis.



Logistic Regression - Coefficients

Ultimately we end up with log(odds of obesity) on the new y-axis...



But how do we fit that line?

We'll see that a bit later...

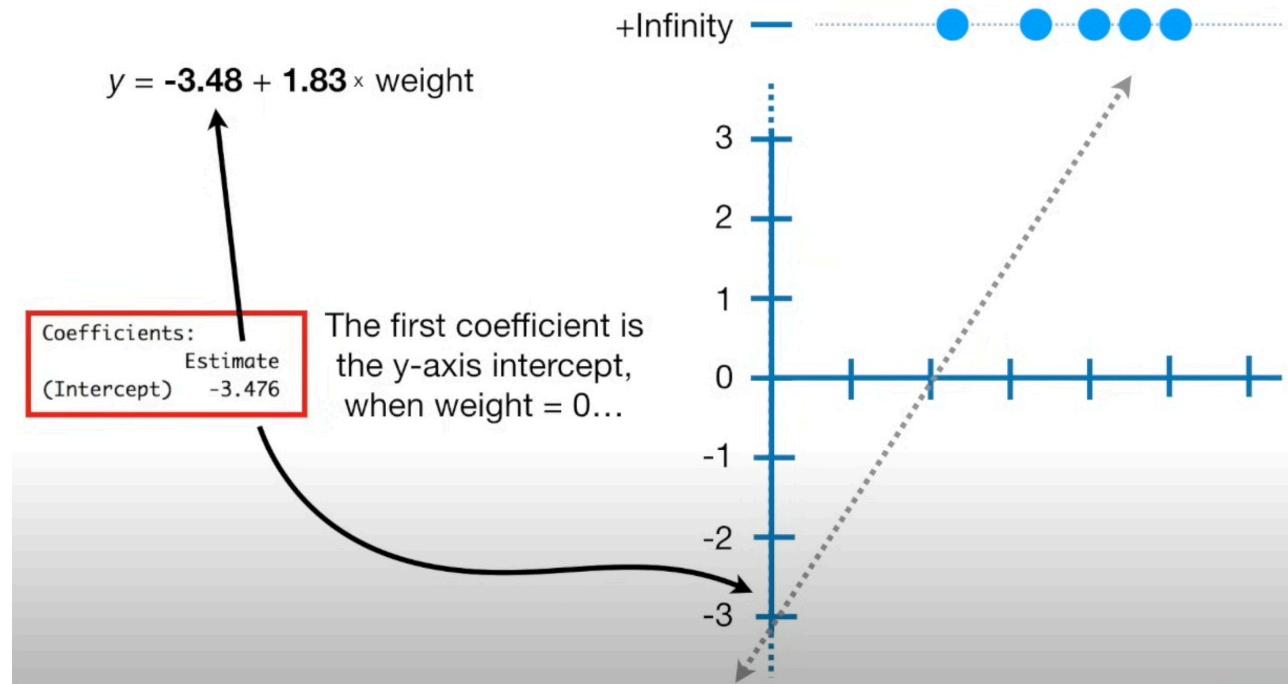
Logistic Regression - Coefficients

The coefficients for the line are what you get when you do logistic regression.

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.476	2.364	-1.471	0.1414
weight	1.825	1.088	1.678	0.0934

Logistic Regression - Coefficients



This means that when weight = 0, the log(odds of obesity) = -3.476

In other words, if you don't weight anything, the odds are against you being obese.

Logistic Regression - Coefficients

Then you have the
Standard Error...

$$y = \mathbf{-3.48 + 1.83 \times weight}$$

...and the z-value is the estimated intercept divided by the standard error.

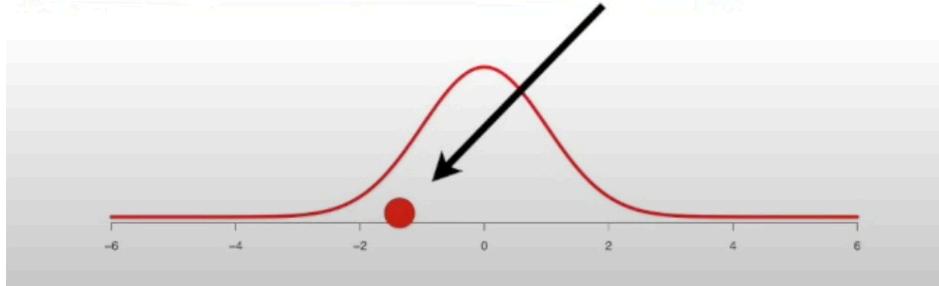
Coefficients:	Estimate	Std. Error	z value
(Intercept)	-3.476	2.364	-1.471

Logistic Regression - Coefficients

In other words, it's the number of standard deviations the estimated intercept is away from 0 on a standard normal curve.

Coefficients:

	Estimate	Std. Error	z value
(Intercept)	-3.476	2.364	-1.471



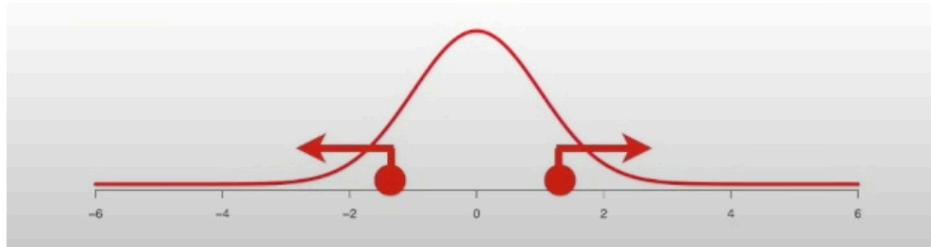
Since the estimate is less than 2 std deviations away from 0, the result is not statistically significant.

Logistic Regression - Coefficients

...and this is confirmed by the large p-value; the area under a standard normal curve that is further than 1.471 standard deviations away from 0.

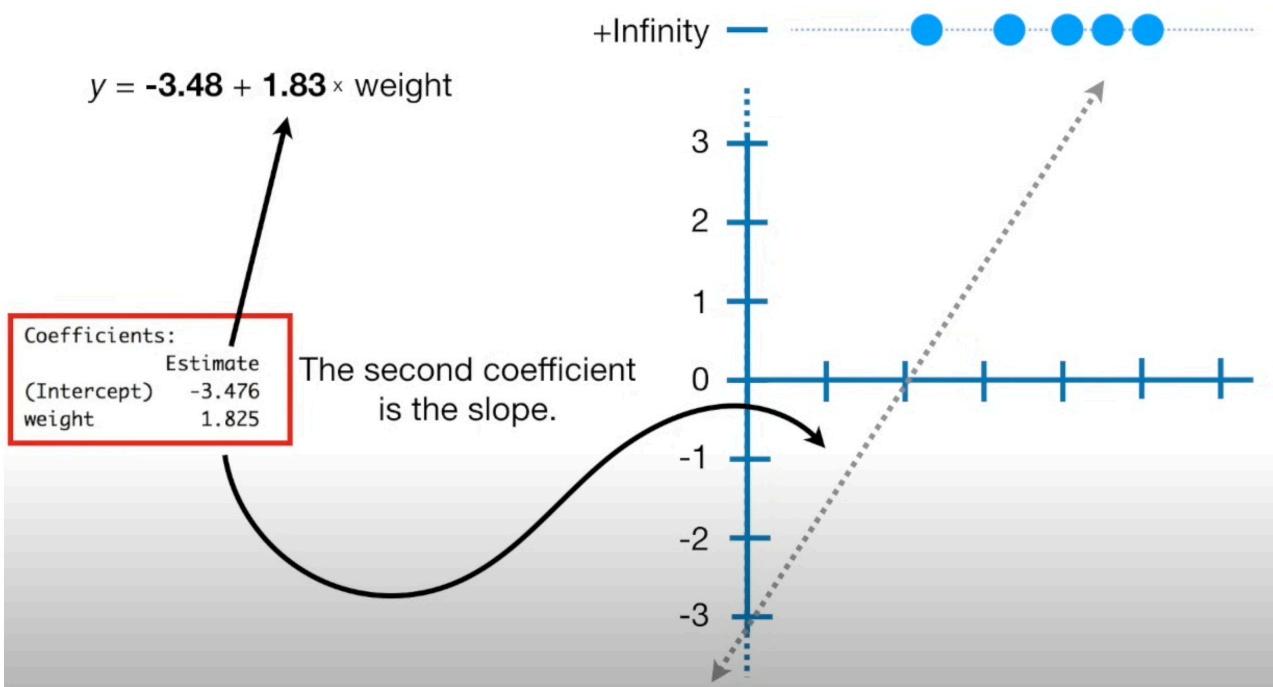
Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-3.476	2.364	-1.471	0.1414



A statistically significant p-value is below 0.05

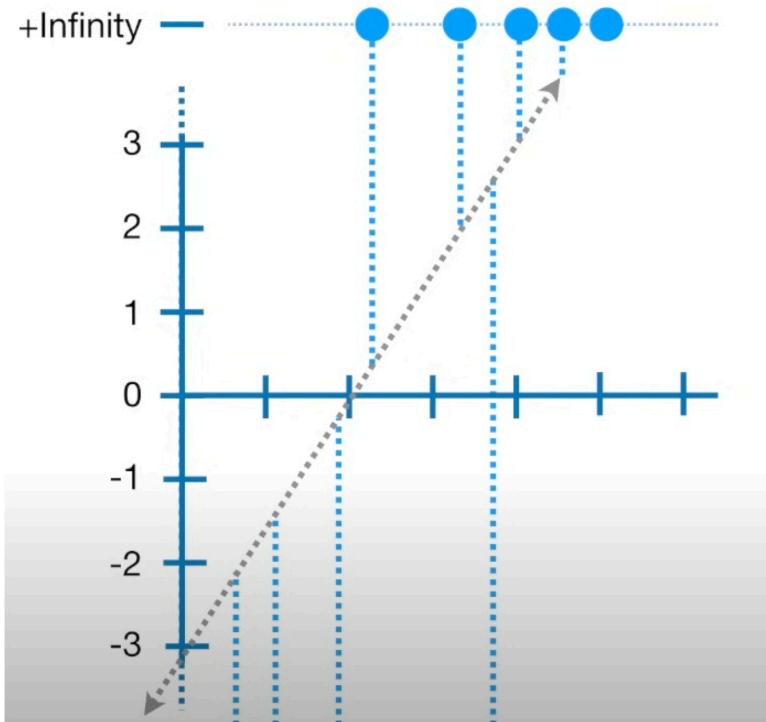
Logistic Regression - Coefficients



This means that for every one unit of weight gained, the $\log(\text{odds of obesity})$ increases by 1.83.

Now that we know what the Coefficients mean, how do we optimize these coefficients to find the best-fitting line?

Logistic Regression - Fitting the line



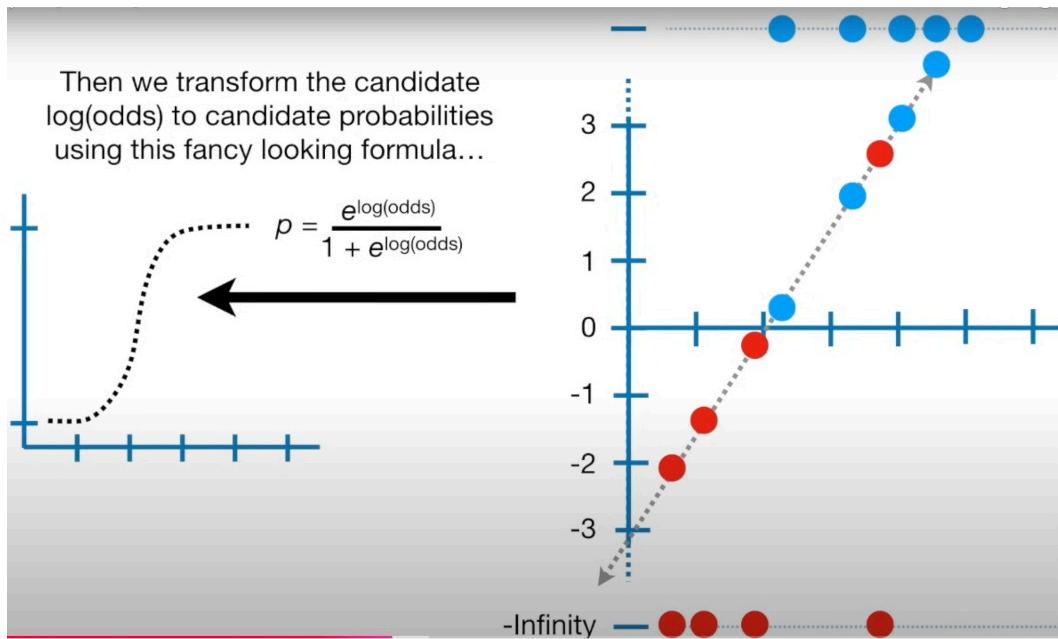
In our new y axis, we go from +inf to -inf...

...and thus we cannot use the least squares method!

We will use the Maximum Likelihood instead

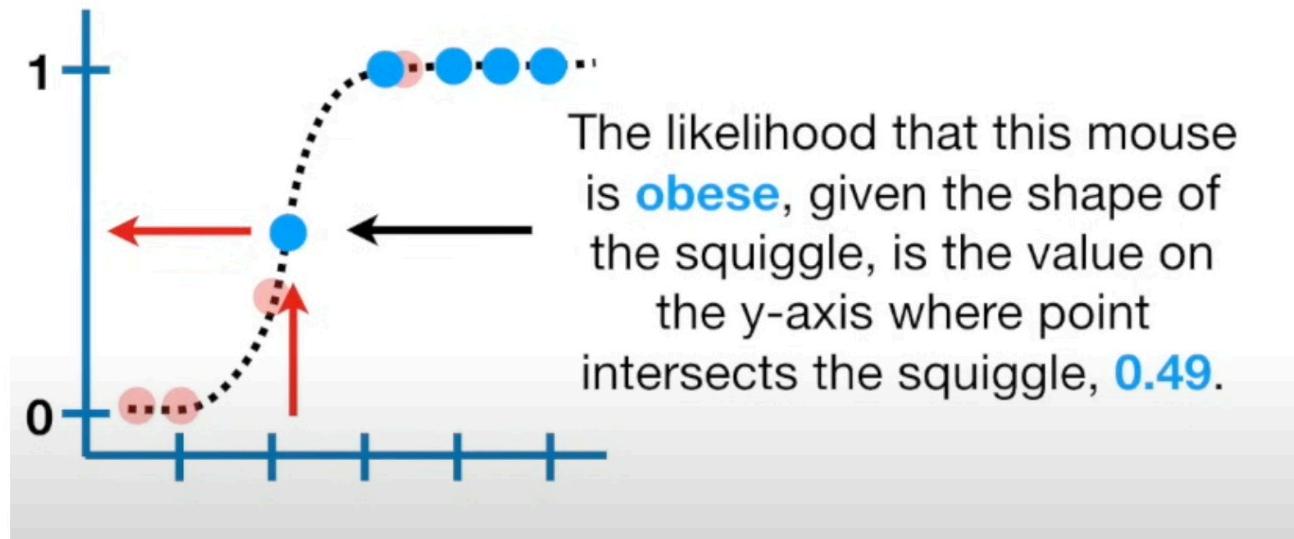
Logistic Regression - Fitting the line

First we project the data points on the line



Logistic Regression - Fitting the line

Then we calculate the likelihood of each data point given the shape of the squiggle...



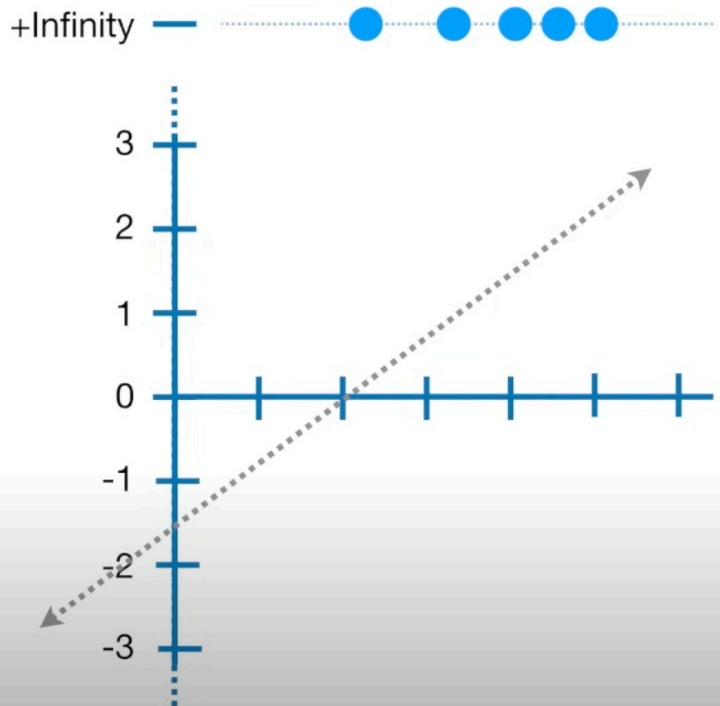
Logistic Regression - Fitting the line

likelihood of data given the squiggle = $0.49 \times 0.9 \times 0.91 \times 0.91 \times 0.92 \times (1 - 0.9) \times (1 - 0.3) \times (1 - 0.01) \times (1 - 0.01)$

...And we sum
the
likelihoods

Logistic Regression - Fitting the line

Now we rotate the line...



...and calculate the likelihoods again
...until we find the line with the maximum likelihood!

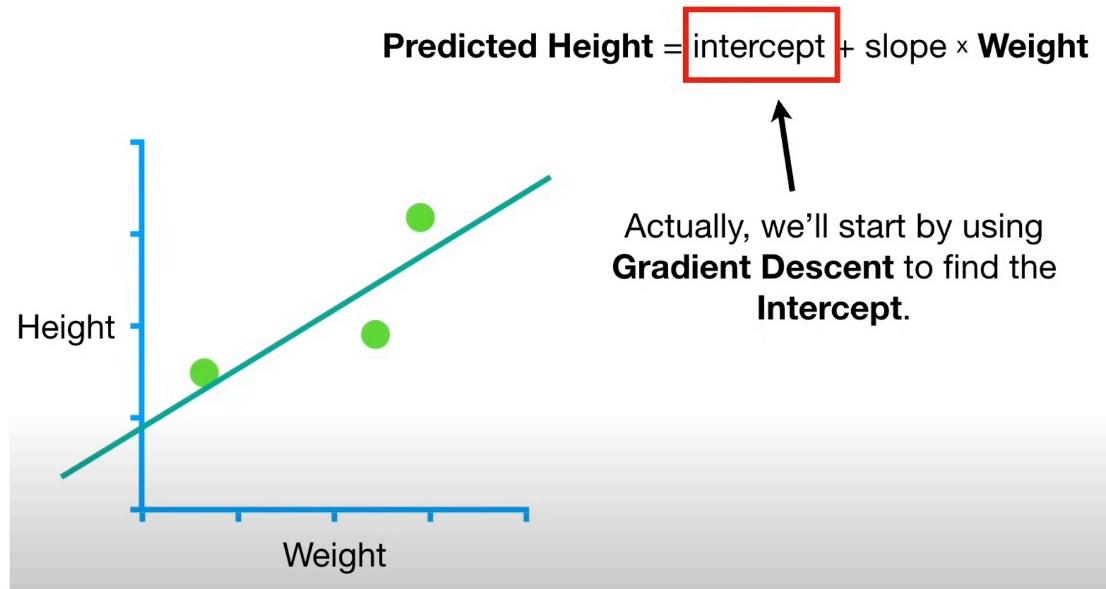
This is great!

But calculating the optimal coefficients for the curve isn't straightforward because there's no simple formula to find them.

Gradient Descent - Intuition

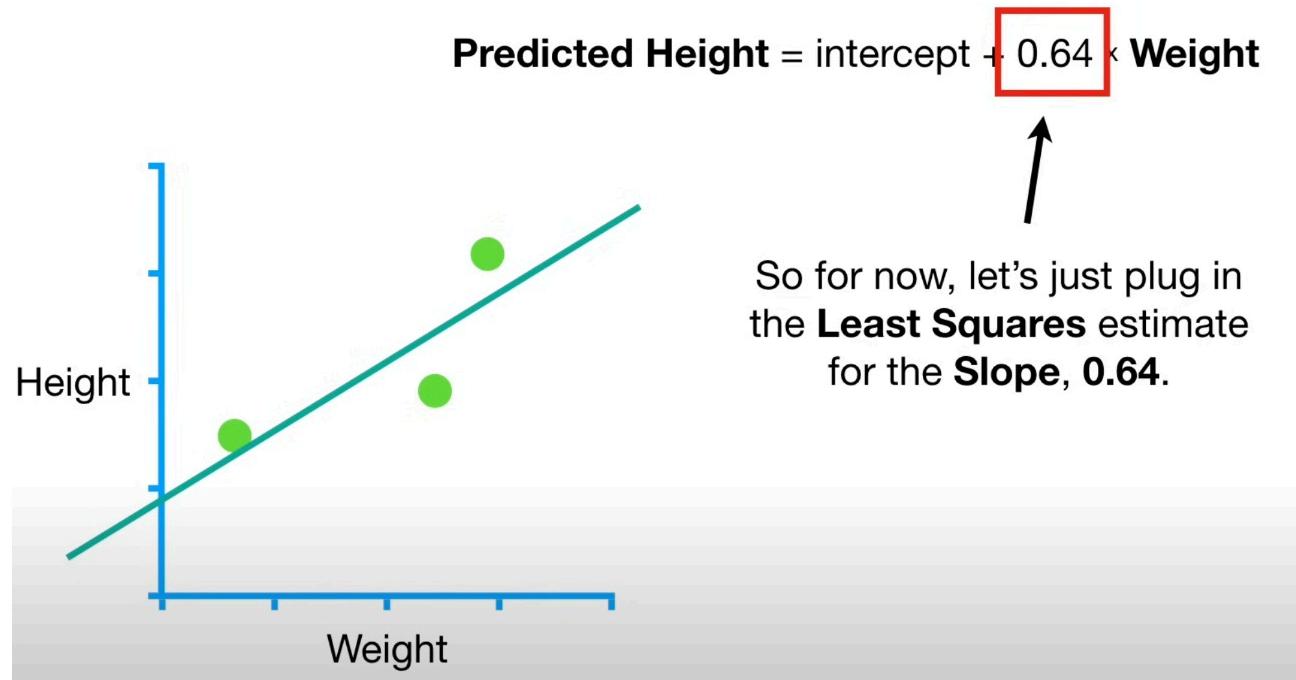
Optimizing the coefficients

This is where
Gradient Descent
come into place



Gradient Descent - Intuition

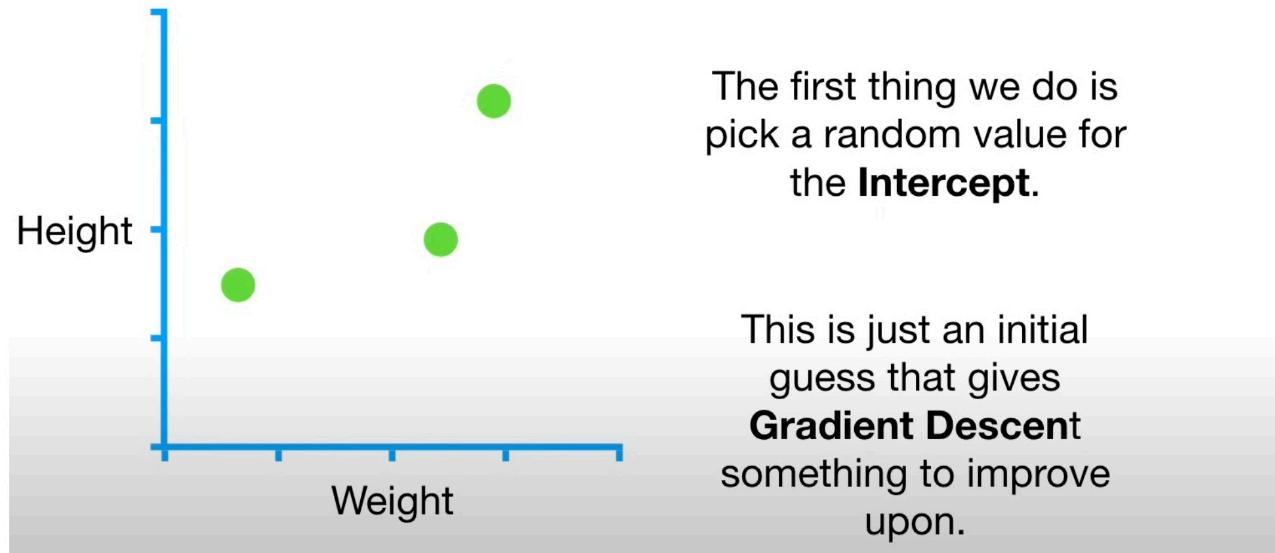
Optimizing the coefficients



Gradient Descent - Intuition

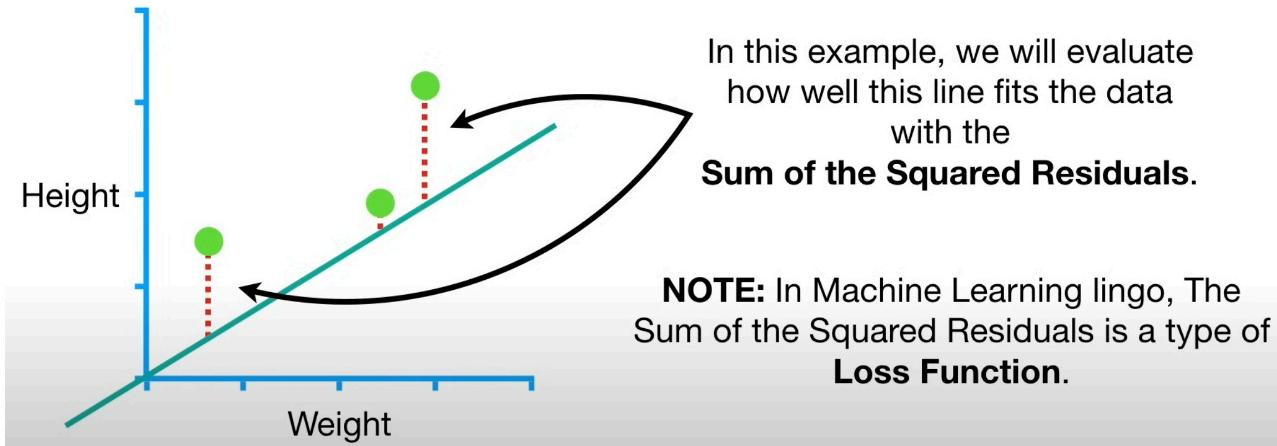
Optimizing the coefficients

$$\text{Predicted Height} = \text{intercept} + 0.64 \times \text{Weight}$$



Gradient Descent - Intuition

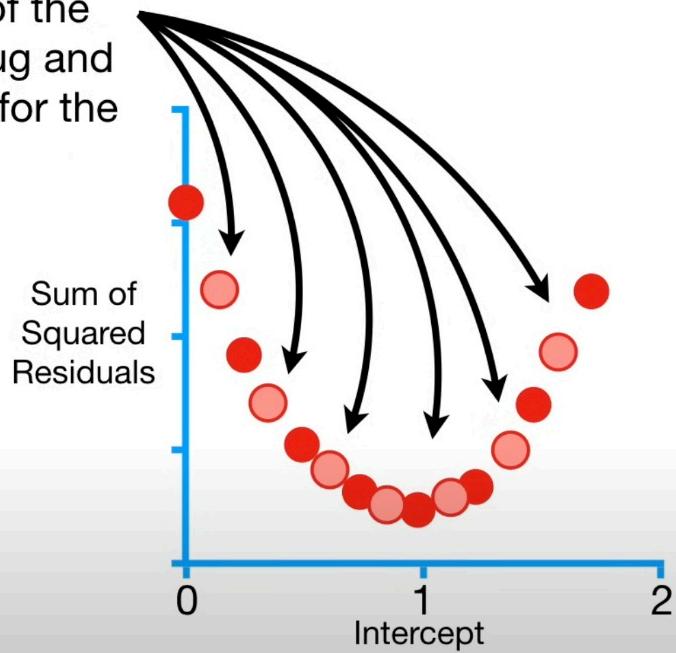
Optimizing the coefficients



Gradient Descent - Intuition

Optimizing the coefficients

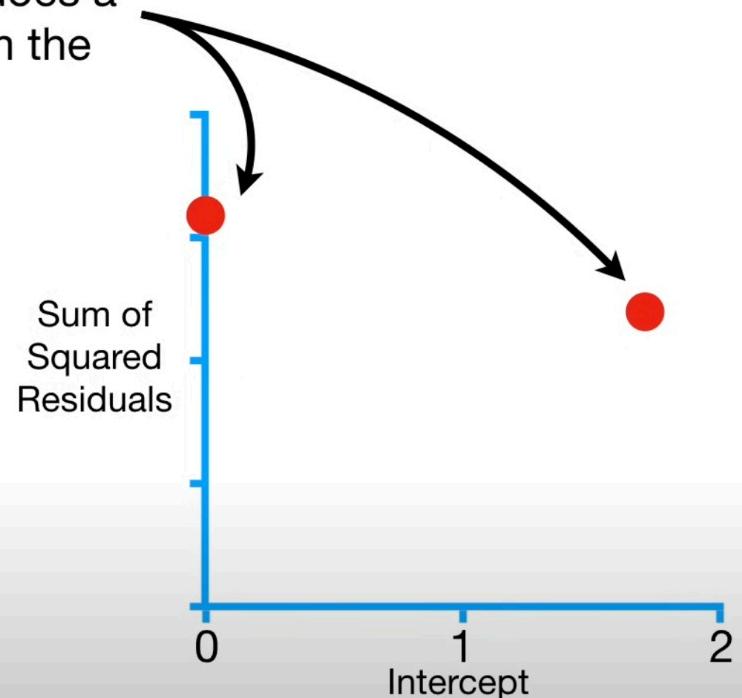
A slow and painful method for finding the minimal Sum of the Squared Residuals is to plug and chug a bunch more values for the **Intercept**.



Gradient Descent - Intuition

Optimizing the coefficients

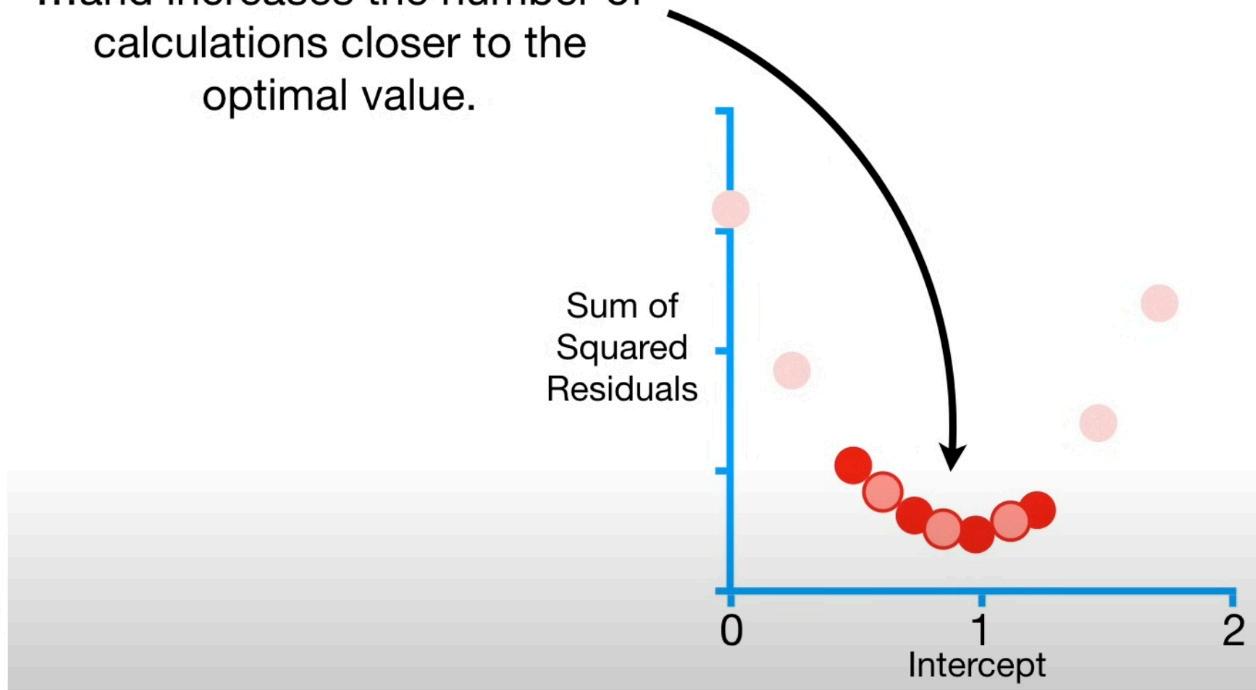
Gradient Descent only does a few calculations far from the optimal solution...



Gradient Descent - Intuition

Optimizing the coefficients

...and increases the number of calculations closer to the optimal value.



Gradient Descent identifies the optimal value by taking big steps when it is far away, and baby steps when it is close

Gradient Descent - Intuition

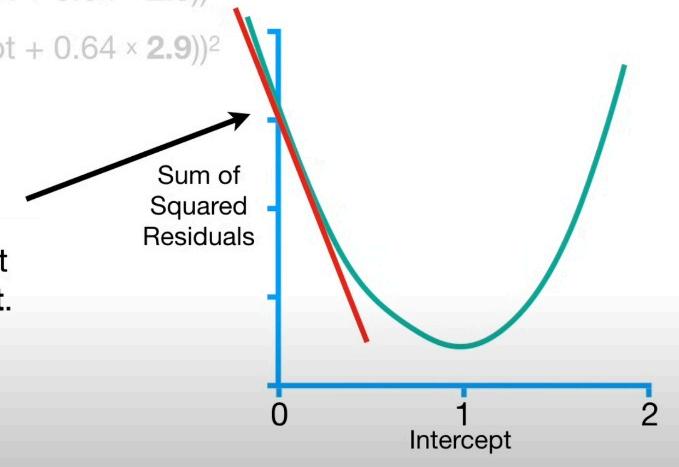
Optimizing the coefficients

How does it work?

Once we have calculated and plotted the sum of squared residuals...

$$\begin{aligned}\text{Sum of squared residuals} = & (1.4 - (\text{intercept} + 0.64 \times 0.5))^2 \\ & + (1.9 - (\text{intercept} + 0.64 \times 2.3))^2 \\ & + (3.2 - (\text{intercept} + 0.64 \times 2.9))^2\end{aligned}$$

...and we can take the derivative of this function and determine the slope at any value for the **Intercept**.

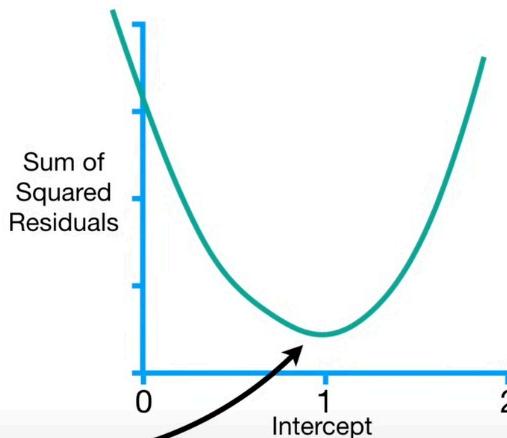


Gradient Descent - Intuition

Optimizing the coefficients

$$\frac{d}{d \text{ intercept}} \text{Sum of squared residuals} =$$
$$-2(1.4 - (\text{intercept} + 0.64 \times 0.5))$$
$$+ -2(1.9 - (\text{intercept} + 0.64 \times 2.3))$$
$$+ -2(3.2 - (\text{intercept} + 0.64 \times 2.9))$$

Now that we have the derivative,
Gradient Descent will use it to find
where the Sum of Squared
Residuals is lowest.



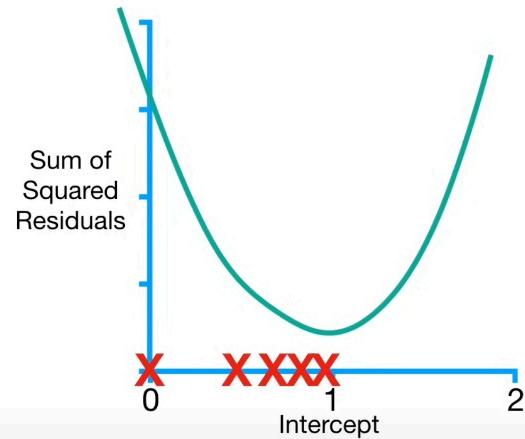
Gradient Descent - Intuition

Optimizing the coefficients

Gradient Descent finds the minimum value by taking steps from an initial guess until it reaches the best value

$$\frac{d}{d \text{ intercept}} \text{Sum of squared residuals} =$$
$$-2(1.4 - (\text{intercept} + 0.64 \times 0.5))$$
$$+ -2(1.9 - (\text{intercept} + 0.64 \times 2.3))$$
$$+ -2(3.2 - (\text{intercept} + 0.64 \times 2.9))$$

This makes **Gradient Descent** very useful when it is not possible to solve for where the derivative = 0, and this is why **Gradient Descent** can be used in so many different situations.



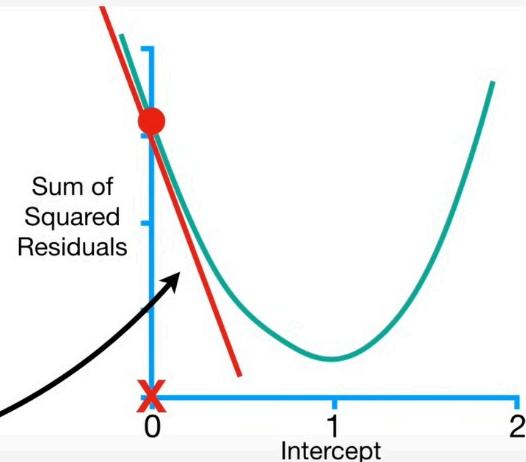
Gradient Descent - Intuition

Optimizing the coefficients

Let's do
Gradient
Descent step
by step

$$\frac{d}{d \text{ intercept}} \text{Sum of squared residuals} =$$
$$-2(1.4 - (0 + 0.64 \times 0.5))$$
$$+ -2(1.9 - (0 + 0.64 \times 2.3))$$
$$+ -2(3.2 - (0 + 0.64 \times 2.9))$$
$$= -5.7$$

So when the **Intercept** = 0,
the slope of the curve = **-5.7**.



Gradient Descent - Intuition

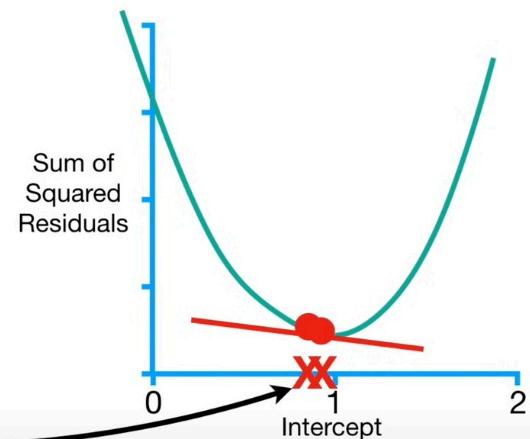
Optimizing the coefficients

The closer we get to the optimal value, the closer the slope of the curve gets to 0.

$$\frac{d}{d \text{ intercept}}$$

$$\begin{aligned} \text{Sum of squared residuals} &= \\ &-2(1.4 - (0 + 0.64 \times 0.5)) \\ &+ -2(1.9 - (0 + 0.64 \times 2.3)) \\ &+ -2(3.2 - (0 + 0.64 \times 2.9)) \\ &= -5.7 \end{aligned}$$

...then we should take baby steps, because we are close to the optimal value...



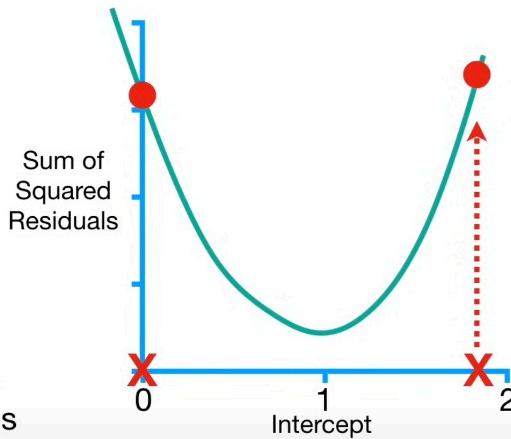
So when we are close to 0...

Gradient Descent - Intuition

Optimizing the coefficients

$$\frac{d}{d \text{ intercept}} \text{Sum of squared residuals} =$$
$$-2(1.4 - (0 + 0.64 \times 0.5))$$
$$+ -2(1.9 - (0 + 0.64 \times 2.3))$$
$$+ -2(3.2 - (0 + 0.64 \times 2.9))$$
$$= -5.7$$

So the size of the step should be related to the slope, since it tells us if we should take a baby step or a big step, but we need to make sure the big step is not too big.



Gradient Descent - Intuition

Optimizing the coefficients

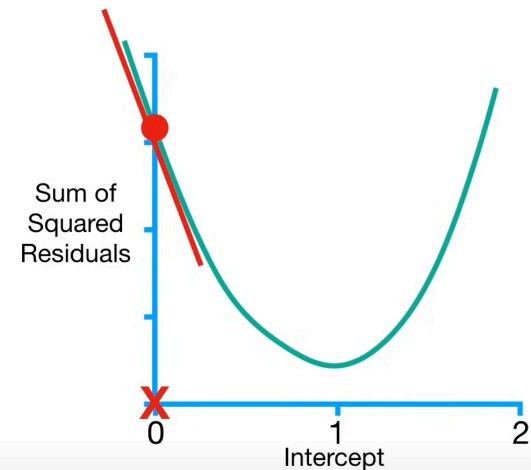
Gradient
Descent
determines
the size of
the step by
multiplying
the slope...

$$\frac{d}{d \text{ intercept}}$$

$$\begin{aligned} \text{Sum of squared residuals} &= \\ &-2(1.4 - (0 + 0.64 \times 0.5)) \\ &+ -2(1.9 - (0 + 0.64 \times 2.3)) \\ &+ -2(3.2 - (0 + 0.64 \times 2.9)) \\ &= -5.7 \end{aligned}$$

Step Size = -5.7×0.1

...by a small number called
The Learning Rate.



Gradient Descent - Intuition

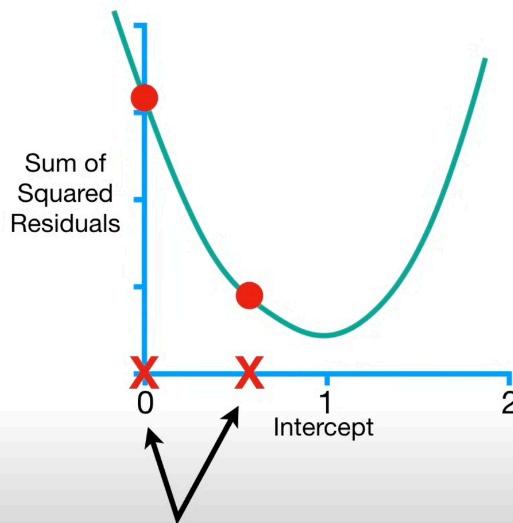
Optimizing the coefficients

$$\frac{d}{d \text{ intercept}} \text{ Sum of squared residuals} =$$
$$-2(1.4 - (0 + 0.64 \times 0.5))$$
$$+ -2(1.9 - (0 + 0.64 \times 2.3))$$
$$+ -2(3.2 - (0 + 0.64 \times 2.9))$$
$$= -5.7$$

$$\text{Step Size} = -5.7 \times 0.1 = -0.57$$

$$\text{New Intercept} = 0 - (-0.57) = 0.57$$

BAM!!!



In one big step, we moved much closer to the optimal value for the **Intercept**.

Gradient Descent - Intuition

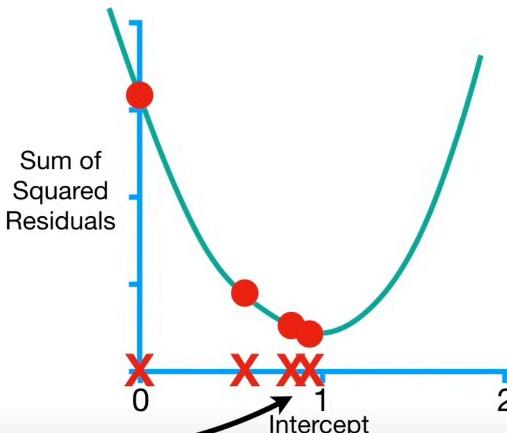
Optimizing the coefficients

$$\frac{d}{d \text{ intercept}} \text{Sum of squared residuals} =$$
$$-2(1.4 - (0.8 + 0.64 \times 0.5))$$
$$+ -2(1.9 - (0.8 + 0.64 \times 2.3))$$
$$+ -2(3.2 - (0.8 + 0.64 \times 2.9))$$
$$= -0.9$$

$$\text{Step Size} = -0.9 \times 0.1 = -0.09$$

$$\text{New Intercept} = 0.8 - (-0.09) = 0.89$$

...and the **New Intercept = 0.89**



Each step gets smaller the closer we get to the optimal value.

We stop when the step size is smaller than 0.001.

Gradient Descent - Intuition

Optimizing the coefficients

To find the optimal values for both intercept and slope, we do the same process, but with both derivatives in respect to intercept and to slope.

NOTE: When you have two or more derivatives of the same function, they are called a **Gradient**.

Gradient Descent - Intuition

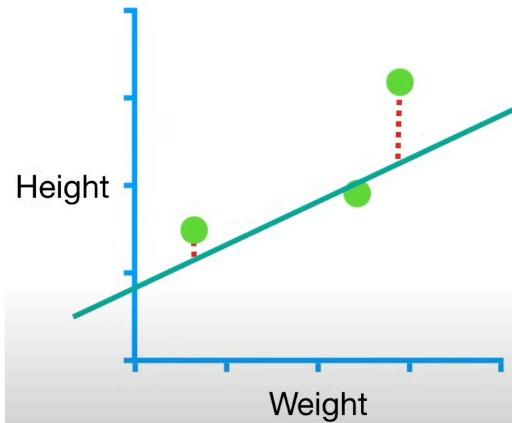
Optimizing the coefficients

We will use this **Gradient** to **descend** to lowest point in the **Loss Function**, which, in this case, is the Sum of the Squared Residuals...

And this is why this method is called **Gradient Descent!**

Gradient Descent - Intuition

Optimizing the coefficients



We now know how **Gradient Descent** optimizes two parameters, the **Slope** and **Intercept**.

If we had more parameters, we would take more derivatives, and everything else stays the same!

Gradient Descent - Intuition

Optimizing the coefficients

Step 1: Take the derivative of the **Loss Function** for each parameter in it.
In fancy Machine Learning Lingo, take the **Gradient** of the **Loss Function**.

Step 2: Pick random values for the parameters.

Step 3: Plug the parameter values into the derivatives (ahem, the **Gradient**).

Step 4: Calculate the Step Sizes: **Step Size = Slope × Learning Rate**

Step 5: Calculate the New Parameters:

New Parameter = Old Parameter - Step Size

Go back to step 3
and repeat until
step is very small
OR you have
reached the
maximum
number of steps

This is it!

Now we have seen 2 different methods to predict outcomes - Linear Regression and Logistic Regression.

But how do we know how good our guesses are?

Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

Chest Pain	Good Blood Circ.	Blocked Arteries	Weight	Heart Disease
No	No	No	125	No
Yes	Yes	Yes	180	Yes
Yes	Yes	No	210	No
...

Imagine we have a dataset of clinical measurements...

Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

Chest Pain	Good Blood Circ.	Blocked Arteries	Weight	Heart Disease
No	No	No	125	No
Yes	Yes	Yes	180	Yes
Yes	Yes	No	210	No
...

... and we want to apply a ML algorithm to predict whether or not someone will develop heart disease

To do this, we could rely on Logistic Regression, but also many others models like Random Forests, K-means, so how do we find the best model for this?

Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

We start by dividing our data into a training dataset and a testing dataset

Then we train all our methods on the training dataset...

...and test each method on the testing dataset

Chest Pain	Good Blood Circ.	Blocked Arteries	Weight	Heart Disease
No	No	No	125	No
...				...

Training Data

Chest Pain	Good Blood Circ.	Blocked Arteries	Weight	Heart Disease
Yes	Yes	No	210	No
...

Testing Data

Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

		Actual	
		Has Heart Disease	Does Not Have Heart Disease
Predicted	Has Heart Disease	True Positives	False Positives
	Does Not Have Heart Disease	False Negatives	True Negatives

We summarize our results into a Confusion Matrix

And now we can calculate useful metrics!

Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

		Actual	
		Has Heart Disease	Does Not Have Heart Disease
Predicted	Has Heart Disease	True Positives	False Positives
	Does Not Have Heart Disease	False Negatives	True Negatives

Sensitivity is the % of patients with heart disease who were correctly identified

$$\text{Sensitivity} = \text{True Positives}/(\text{True Positives} + \text{False Negatives})$$

Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

		Actual	
		Has Heart Disease	Does Not Have Heart Disease
Predicted	Has Heart Disease	True Positives	False Positives
	Does Not Have Heart Disease	False Negatives	True Negatives

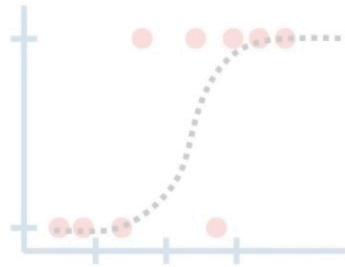
Specificity is the % of patients without heart disease who were correctly identified

$$\text{Specificity} = \text{True Negatives} / (\text{True Negatives} + \text{False Positives})$$

Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

Let's start by calculating **Sensitivity** for this **Logistic Regression**.



$$\begin{aligned}\text{Sensitivity} &= \\ 139 &/ 171 = \\ 0.81\end{aligned}$$

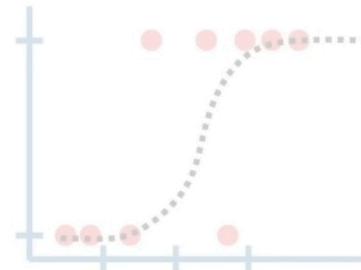
	Has Heart Disease	Does Not Have Heart Disease	
Has Heart Disease	139	20	
Does Not Have Heart Disease	32	112	

81% of patients with heart disease where correctly identified with this model.

Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

Now let's calculate the **Specificity**...



$$\text{Specificity} = \\ 20/132 = \\ 0.85$$

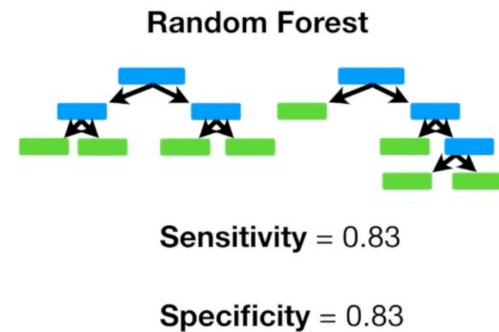
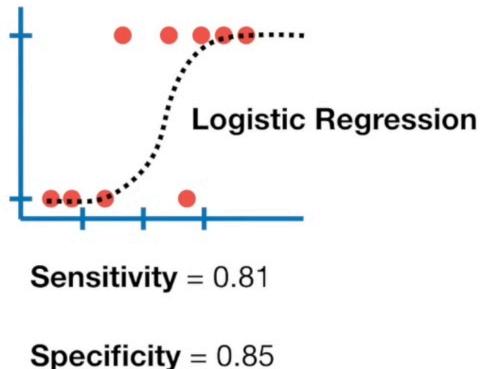
		Has Heart Disease	Does Not Have Heart Disease
Has Heart Disease	Has Heart Disease	139	20
	Does Not Have Heart Disease	32	112

85% of patients without heart disease were correctly identified with this model.

Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

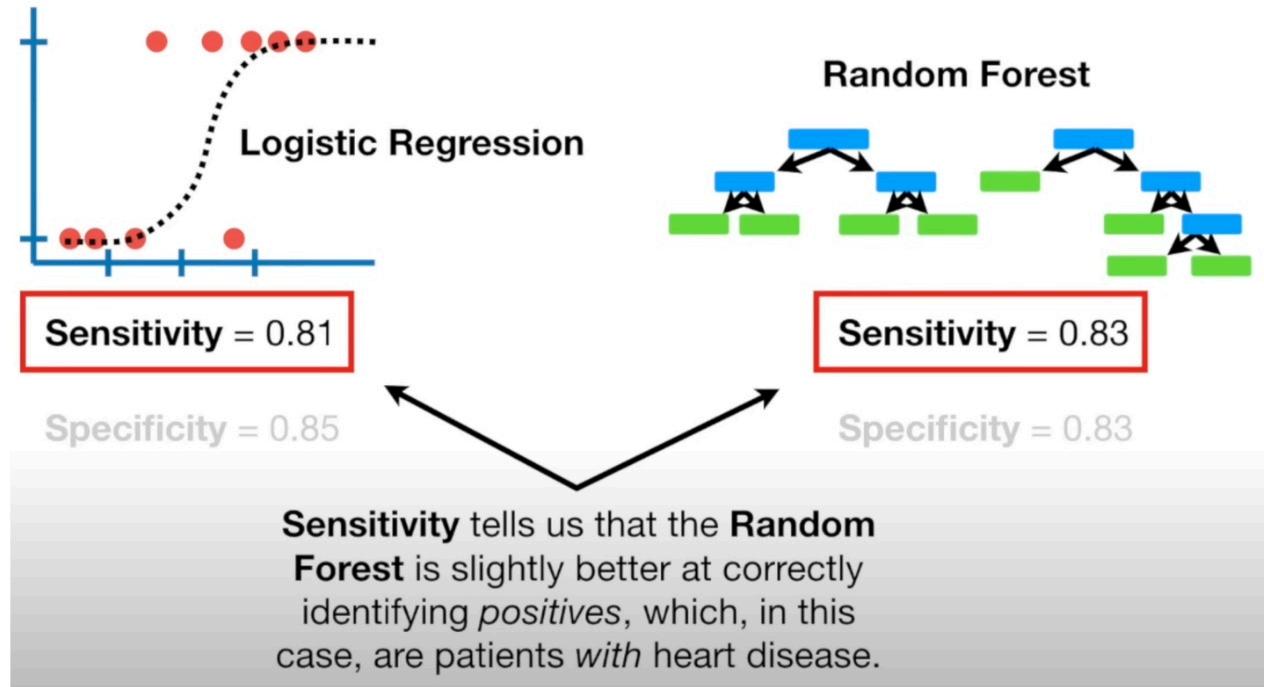
Now let's compare
these values...



...to the values we calculated for the
Random Forest.

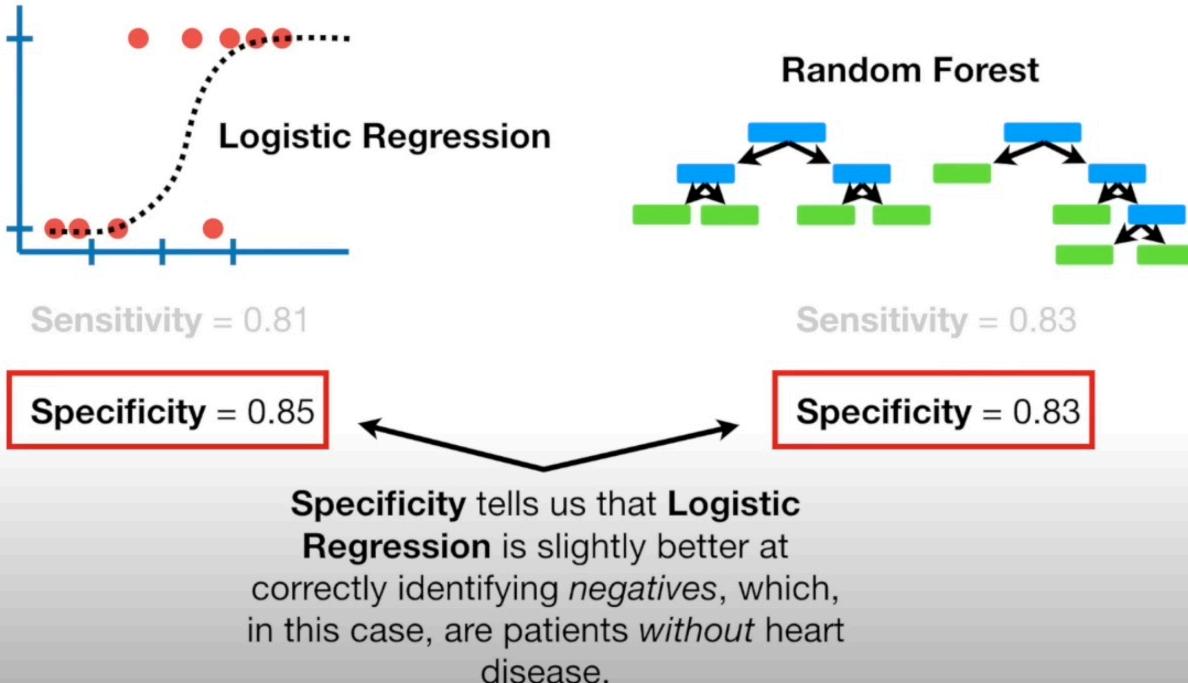
Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC



Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC



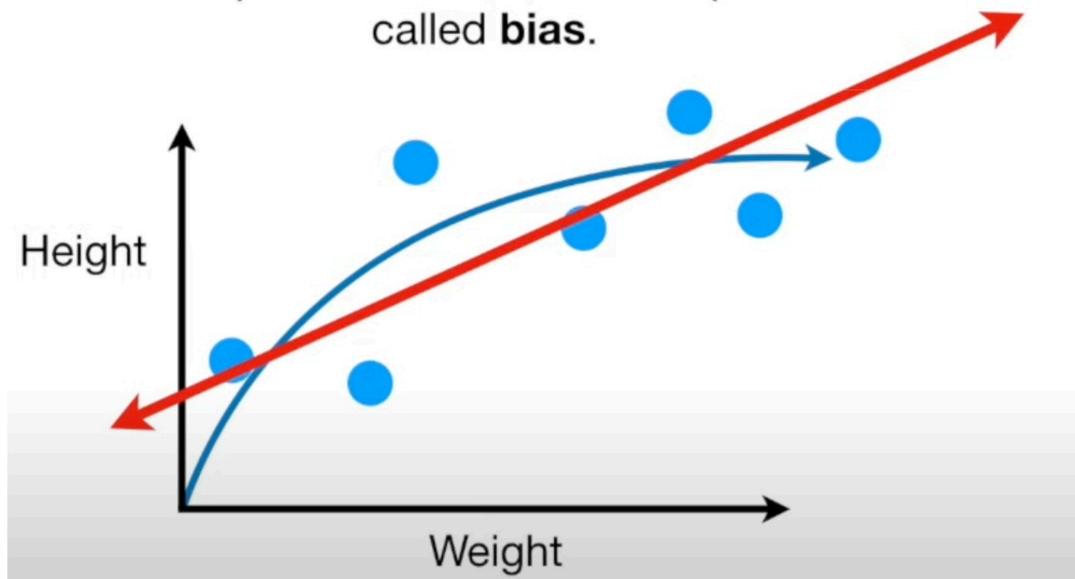
We will choose the Random Forest if correctly identifying people with heart disease is more important than correctly identifying people without heart disease.

Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

Because linear regression produces a straight line, the model will never capture a perfect relationship

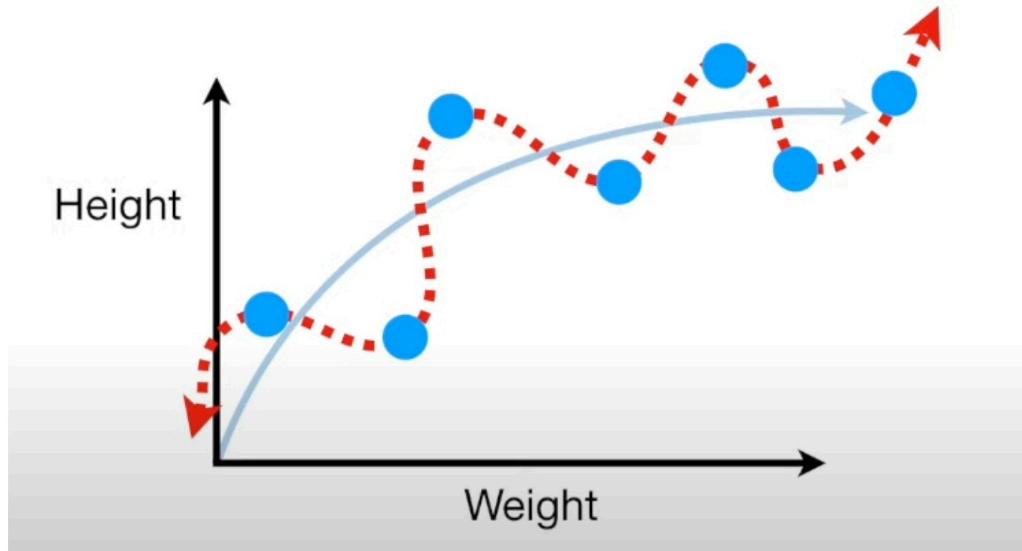
The inability for a machine learning method (like linear regression) to capture the true relationship is called **bias**.



Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

Because the **Squiggly Line** can handle the arc in the true relationship between weight and height, it has very little **bias**.



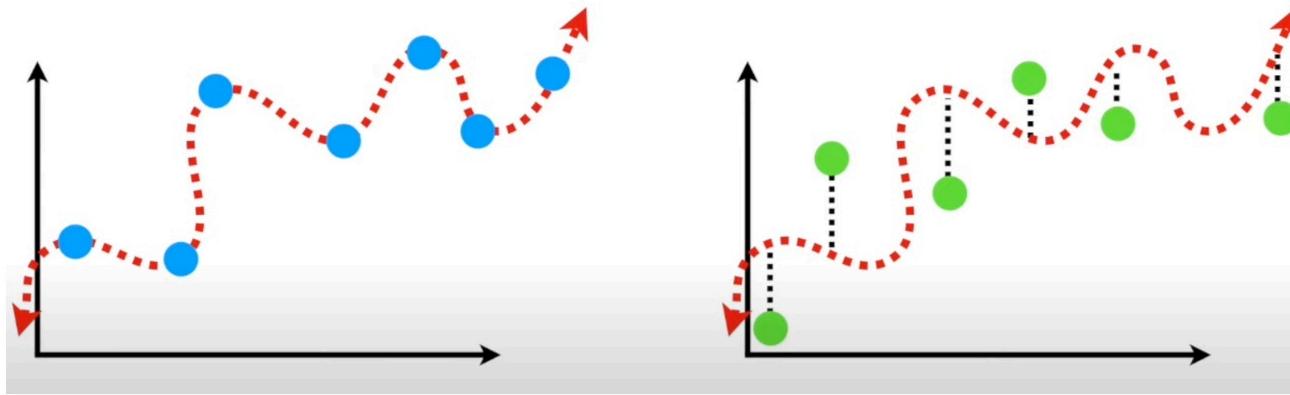
It fits perfectly this dataset...

...but it may not fit very well on other datasets

Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

In Machine Learning lingo, the difference in fits between data sets is called **Variance**.

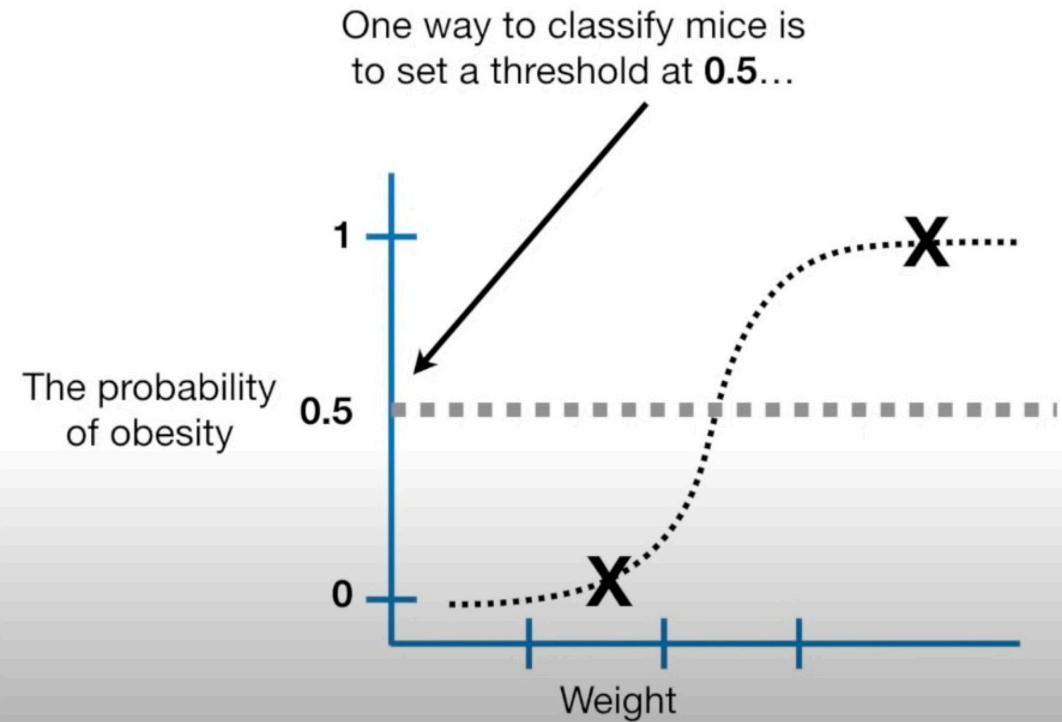


When a model has low bias but high variance, we call it **Overfitting**

Evaluation metrics

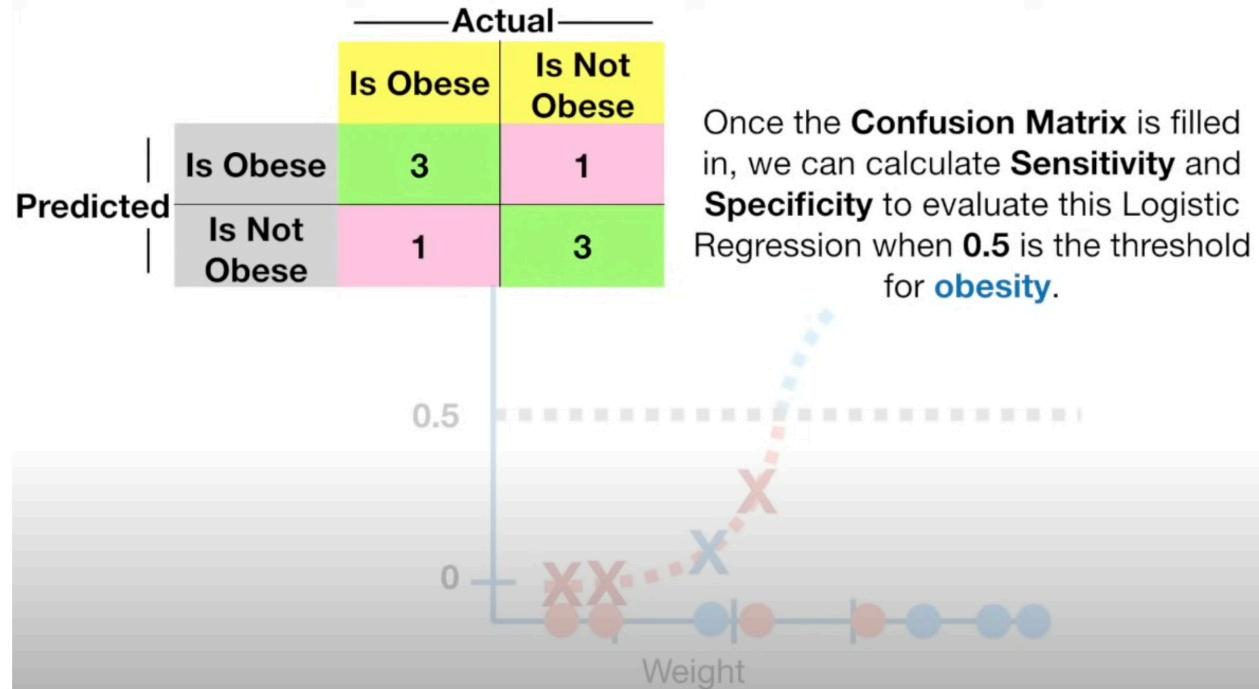
Confusion Matrix, Sensitivity, Specificity, ROC, AUC

Now let's talk about
Logistic Regression
again



Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC



But how do we know this is the right threshold?

Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

If we made 1 confusion matrix per threshold, it would result in a lot of confusion matrices...

	Is Obese	Is Not Obese		Is Obese	Is Not Obese		Is Obese	Is Not Obese
Is Obese	4	0	Is Obese	4	0	Is Obese	3	1
Is Not Obese	0	0	Is Not Obese	0	0	Is Not Obese	1	3
	Is Obese	Is Not Obese		Is Obese	Is Not Obese		Is Obese	Is Not Obese
Is Obese	4	0	Is Obese	3	0	Is Obese	4	0
Is Not Obese	0	1	Is Not Obese	1	0	Is Not Obese	0	4

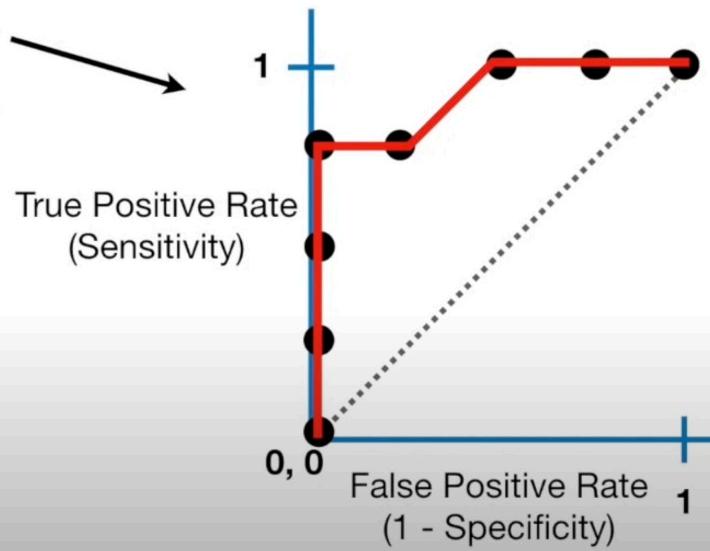
Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

So instead of being overwhelmed
with confusion matrices,

**Receiver Operator Characteristic
(ROC) graphs**

provide a simple way to summarize
all of the information.

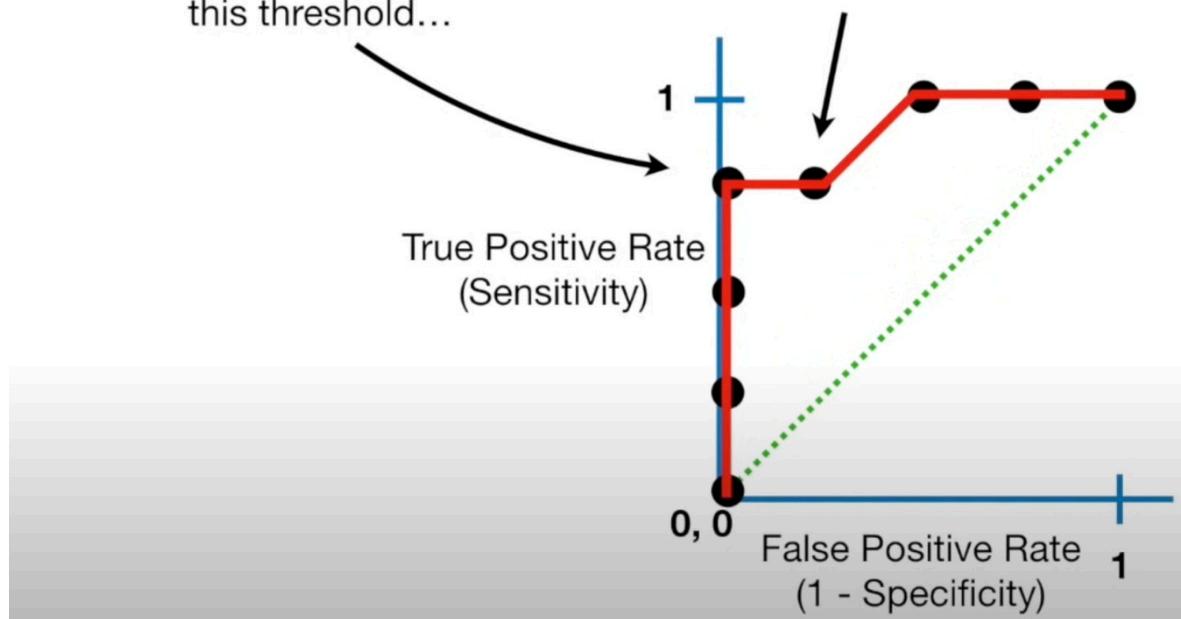


Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

Without having to sort through the confusion matrices, I can tell that this threshold...

...is better than this threshold.

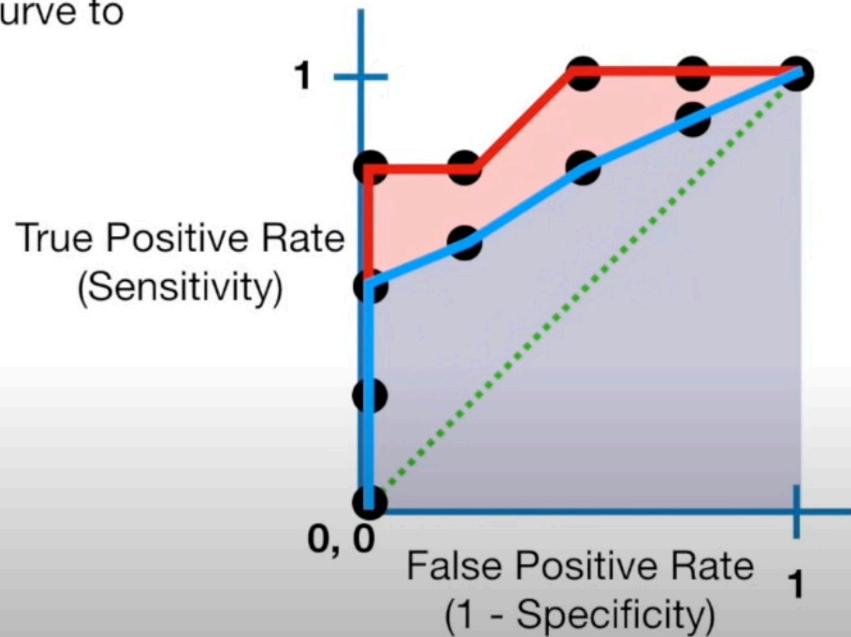


Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

Now that we know what ROC is, let's talk about the Area Under the Curve

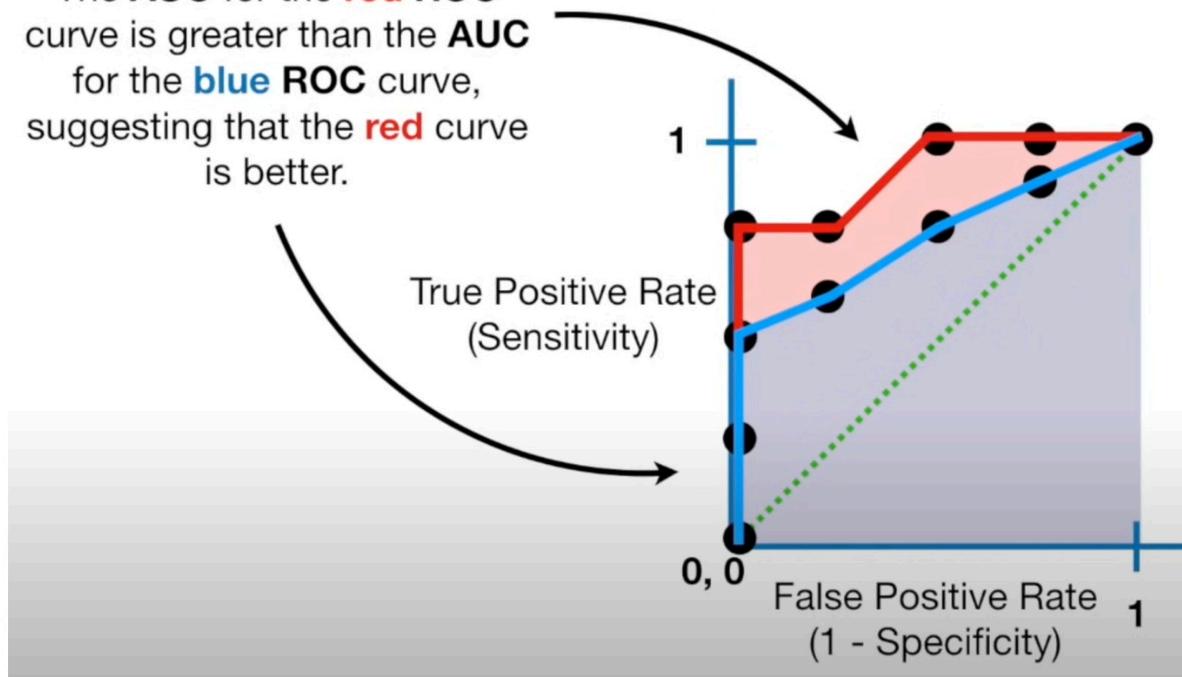
The **AUC** makes it easy to compare one **ROC** curve to another.



Evaluation metrics

Confusion Matrix, Sensitivity, Specificity, ROC, AUC

The **AUC** for the **red ROC** curve is greater than the **AUC** for the **blue ROC** curve, suggesting that the **red** curve is better.



So if the red ROC curve represented Logistic Regression and the blue one Random Forest, you would choose Logistic Regression.